



FORWARD PRIVATE SEARCHABLE ENCRYPTION

DATE

27/10/2016

ACM CCS - RAPHAEL BOST

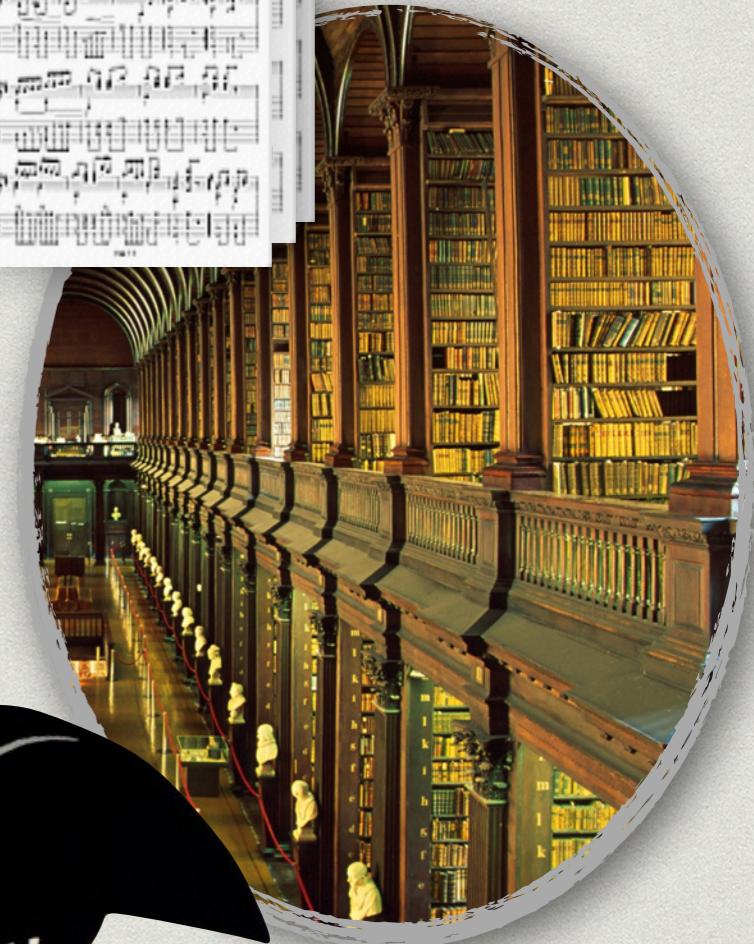


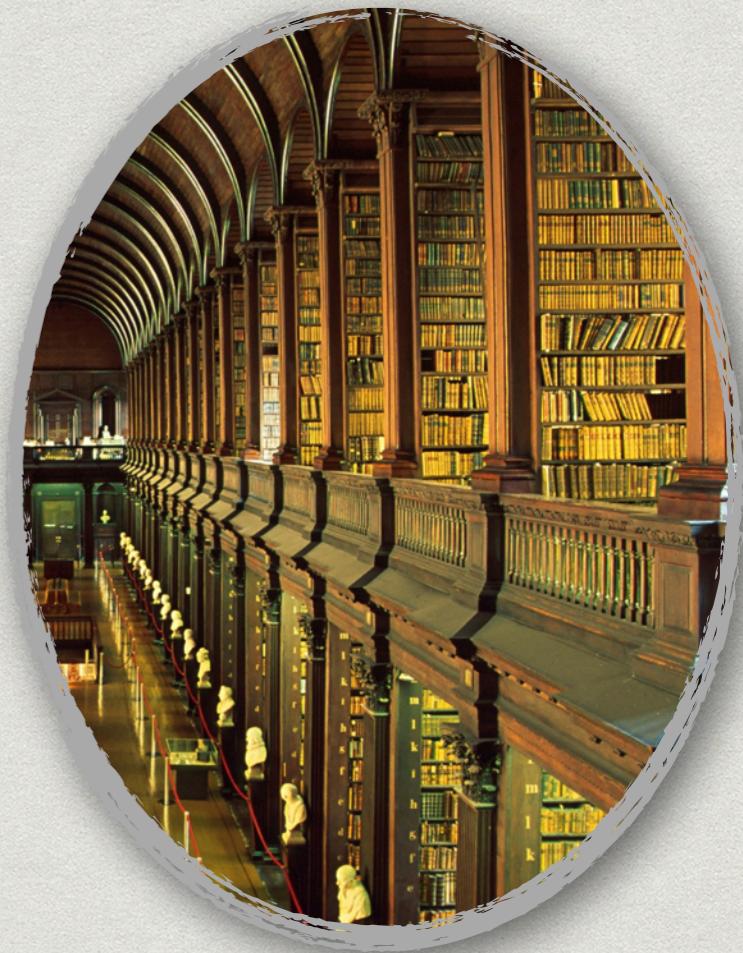


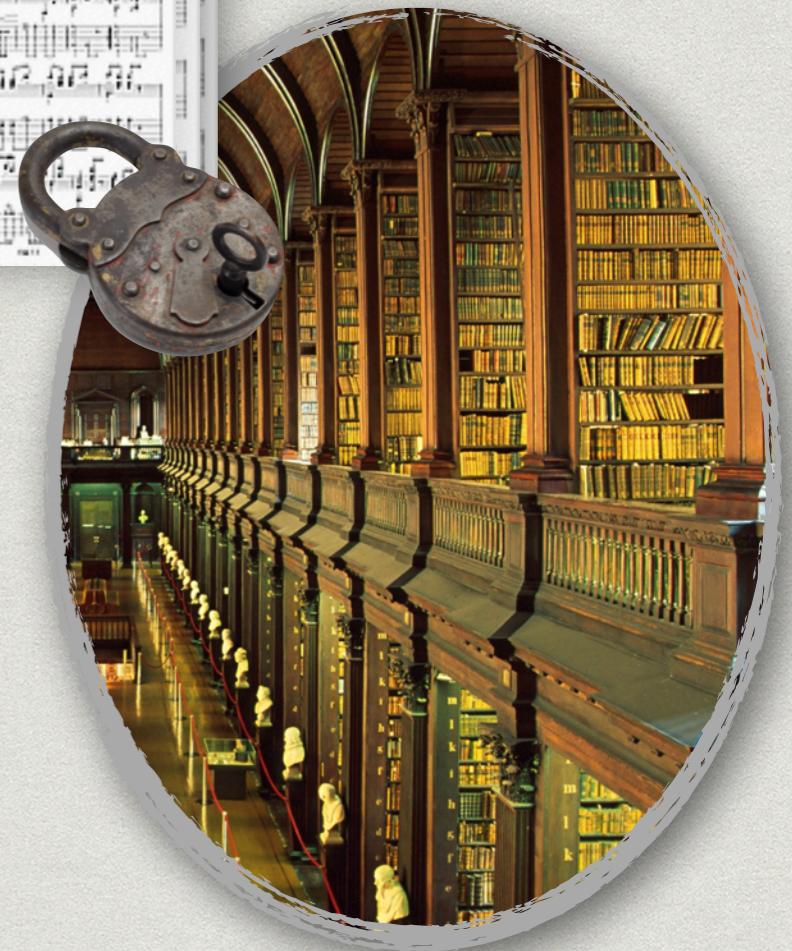


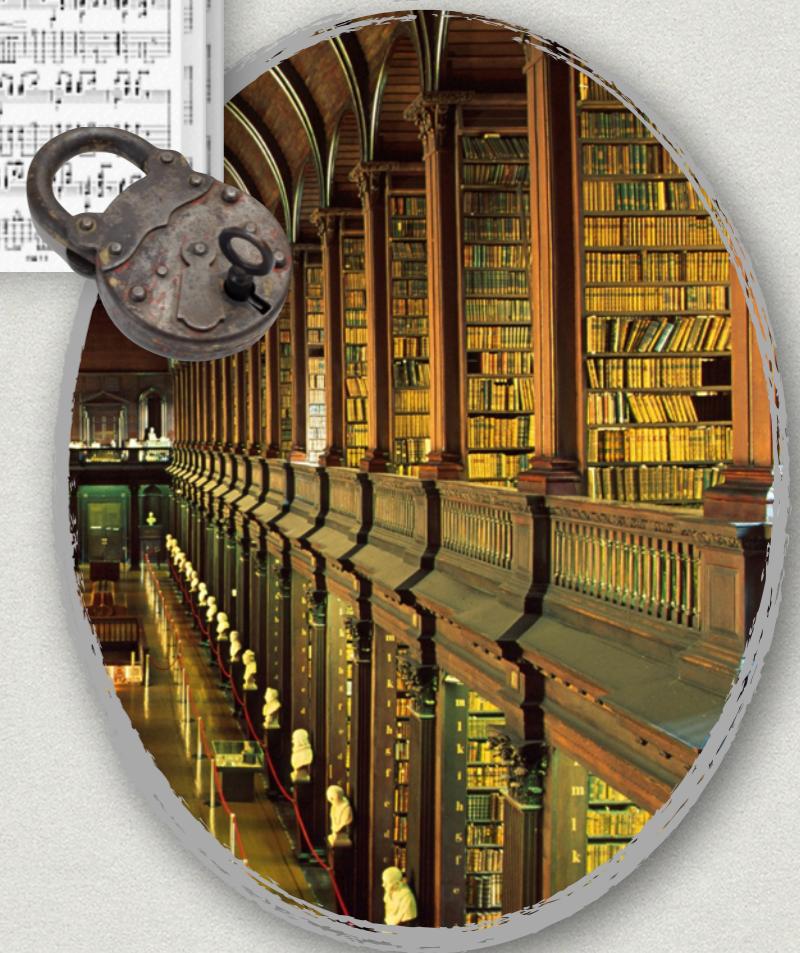


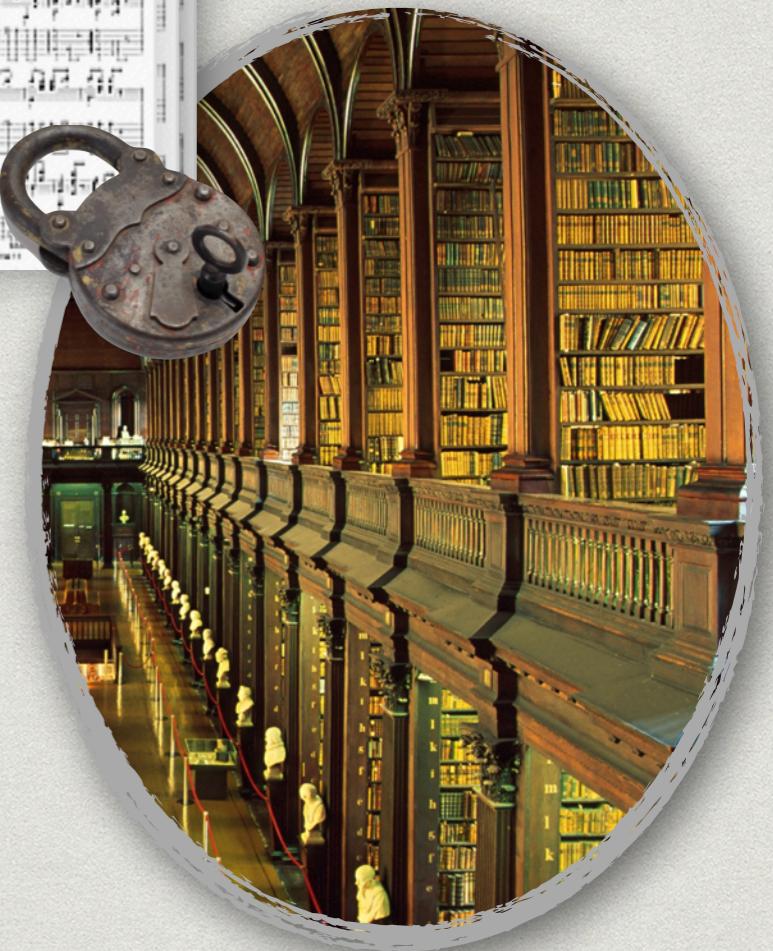


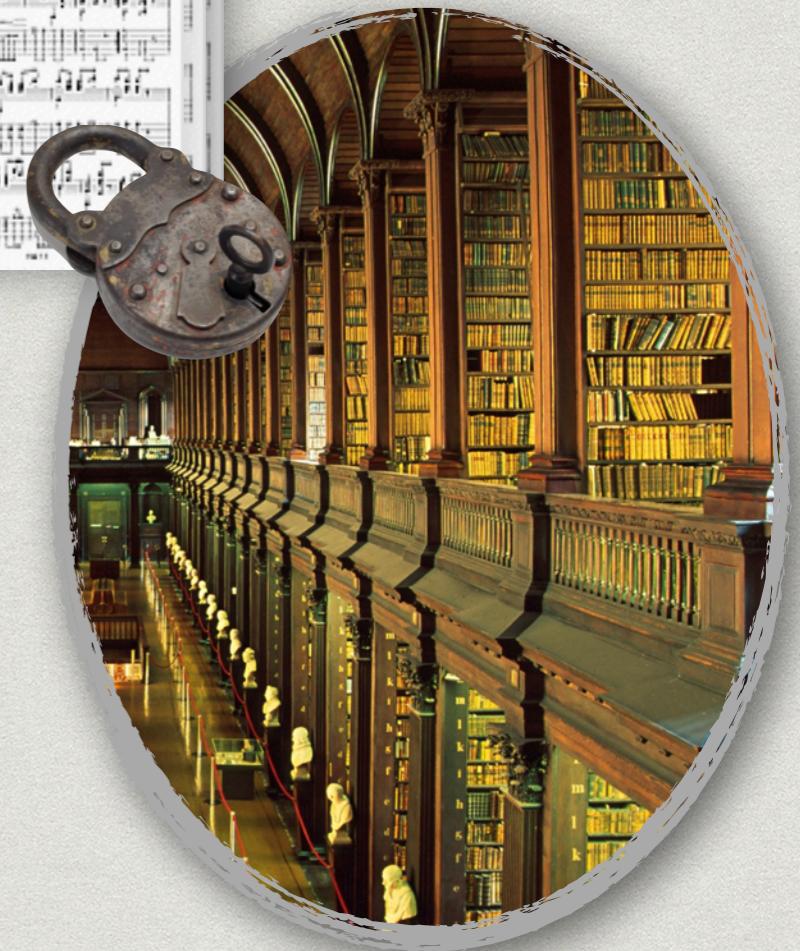


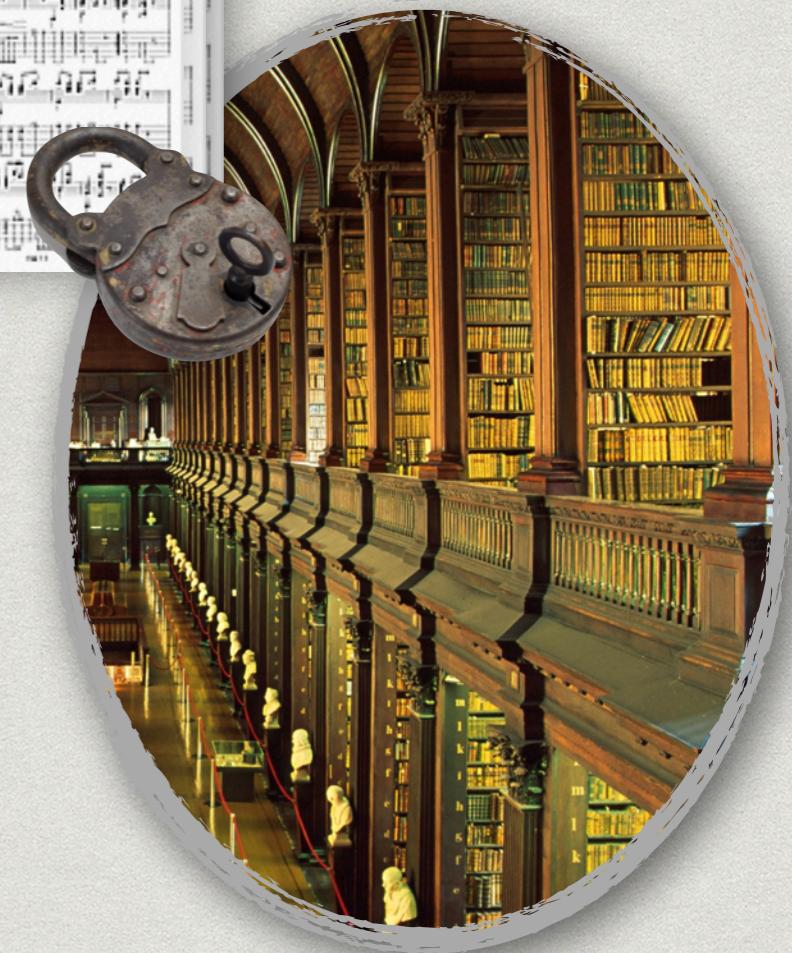












Searchable Encryption

Outsource data

Searchable Encryption

Outsource data

- * securely

Searchable Encryption

Outsource data

- * securely
- * keep search functionalities

Searchable Encryption

Outsource data

- * securely
- * keep search functionalities
- * aimed at efficiency

Generic Solutions

Generic Solutions

Fully Homomorphic Encryption, MPC,
ORAM

- ✓ Perfect security
- ✗ Large overhead (computation,
communication)

Ad-hoc Constructions

Ad-hoc Constructions

Can we get more efficient solutions?

Ad-hoc Constructions

Can we get more efficient solutions?

- * Yes, but ...

Ad-hoc Constructions

Can we get more efficient solutions?

- * Yes, but ...
- * ... we have to leak some information

Ad-hoc Constructions

Can we get more efficient solutions?

- * Yes, but ...
- * ... we have to leak some information

Security/performance tradeoff

Property Preserving Encryption

Property Preserving Encryption

Deterministic Encryption, OPE, ORE

Property Preserving Encryption

Deterministic Encryption, OPE, ORE

✓ Legacy compatible

Property Preserving Encryption

Deterministic Encryption, OPE, ORE

- ✓ Legacy compatible

- ✓ Very Efficient

Property Preserving Encryption

Deterministic Encryption, OPE, ORE

- ✓ Legacy compatible
- ✓ Very Efficient
- ✗ Not secure in practice (cf. next talk)

Index-Based SE [CGKO'06]

Index-Based SE [CGKO'06]

Encryption of the reversed index

Index-Based SE [CGKO'06]

Encryption of the reversed index

- * Search leakage :
- * repetition of queries

Index-Based SE [CGKO'06]

Encryption of the reversed index

- * Search leakage :
 - * repetition of queries
- * Update leakage:
 - * updated documents
 - * repetition of updated keywords

File Injection Attacks [ZKP'16]

Non-adaptive file injection attacks

- * Insert purposely crafted documents in the DB.
Use binary search to recover the query

D ₁	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₂	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₃	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈

File Injection Attacks [ZKP'16]

Non-adaptive file injection attacks

- * Insert purposely crafted documents in the DB.
Use binary search to recover the query

D ₁	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₂	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₃	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈

File Injection Attacks [ZKP'16]

Non-adaptive file injection attacks

- * Insert purposely crafted documents in the DB.
Use binary search to recover the query

D ₁	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₂	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₃	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈

The diagram illustrates a database structure with three rows of documents (D₁, D₂, D₃) and eight keys (k₁ to k₈). The first two columns are highlighted with orange boxes. The third column, containing k₃, is also highlighted with an orange box. An orange arrow points upwards from the bottom row towards the highlighted area, indicating a search operation.

File Injection Attacks [ZKP'16]

Non-adaptive file injection attacks

- * Insert purposely crafted documents in the DB.
Use binary search to recover the query

D ₁	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₂	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈
D ₃	k ₁	k ₂	k ₃	k ₄	k ₅	k ₆	k ₇	k ₈

log K injected documents

File Injection Attacks [ZKP'16]

- * Insert purposely crafted documents.
Use binary search to recover the query

→ Counter measure:
no more than T kw./doc.

$(K/T) \cdot \log T$ injected documents

- * Adaptive version of the attack

$(K/T) + \log T$ injected documents

File Injection Attacks [ZKP'16]

- * Insert purposely crafted documents.
Use binary search to recover the query

→ Counter measure:
no more than T kw./doc.

$(K/T) \cdot \log T$ injected documents

- * Adaptive version of the attack

$(K/T) + \log T$ injected documents

$\log T$ with prior knowledge

Adaptive File Injection

- * The adaptive attack uses the update leakage:

Adaptive File Injection

- * The adaptive attack uses the update leakage:
- * Most SE schemes leak if a newly inserted document matches a previous query

Adaptive File Injection

- * The adaptive attack uses the update leakage:
- * Most SE schemes leak if a newly inserted document matches a previous query
- * We need SE schemes with oblivious updates

Adaptive File Injection

- * The adaptive attack uses the update leakage:
- * Most SE schemes leak if a newly inserted document matches a previous query
- * We need SE schemes with oblivious updates

Forward Privacy

Forward Privacy

- * Forward private: an update does not leak any information on the updated keywords

Forward Privacy

- * Forward private: an update does not leak any information on the updated keywords
- * Secure online build of the EDB

Forward Privacy

- * Forward private: an update does not leak any information on the updated keywords
- * Secure online build of the EDB
- * Only one existing scheme so far [SPS'14]

Forward Privacy

- * Forward private: an update does not leak any information on the updated keywords
- * Secure online build of the EDB
- * Only one existing scheme so far [SPS'14]
 - ORAM-like construction

Forward Privacy

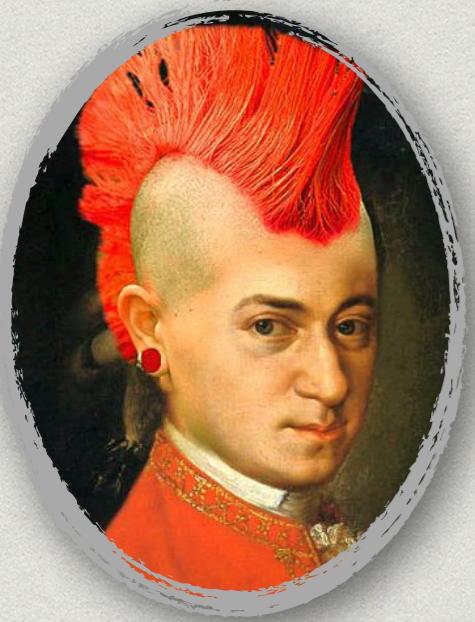
- * Forward private: an update does not leak any information on the updated keywords
- * Secure online build of the EDB
- * Only one existing scheme so far [SPS'14]
 - ORAM-like construction
 - ✗ Inefficient updates

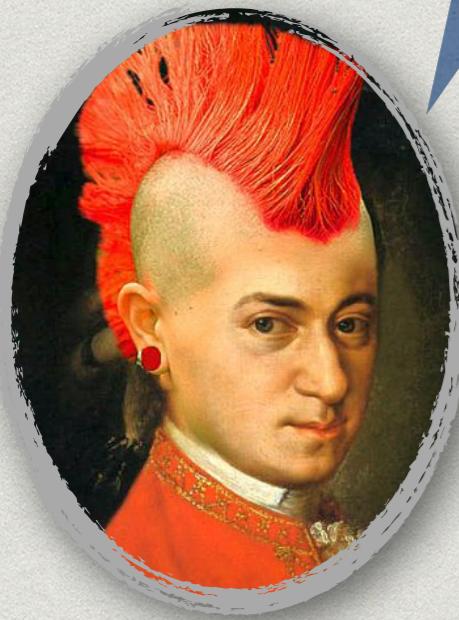
Forward Privacy

- * Forward private: an update does not leak any information on the updated keywords
- * Secure online build of the EDB
- * Only one existing scheme so far [SPS'14]
 - ORAM-like construction
 - ✗ Inefficient updates
 - ✗ Large client storage

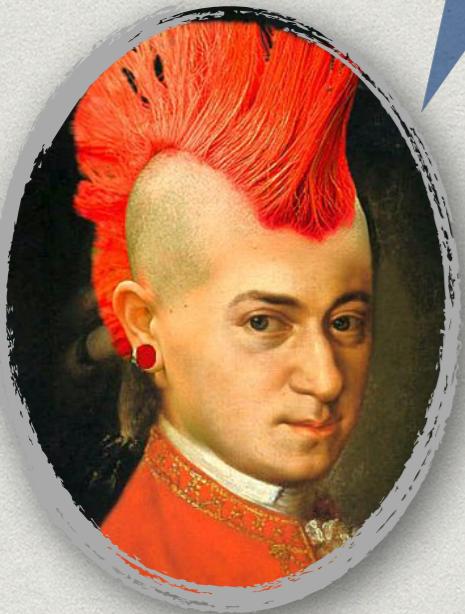
Σοφος

- * Forward private index-based scheme
- * Low search and update overhead
- * Simpler than [SPS'14]





Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w



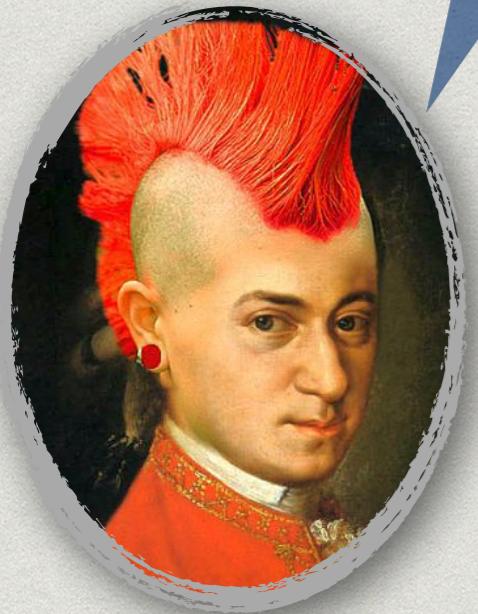
Add (ind_1, \dots, ind_c) to w

$UT_1(w)$

$UT_2(w)$

....

$UT_c(w)$



Add (ind_1, \dots, ind_c) to w

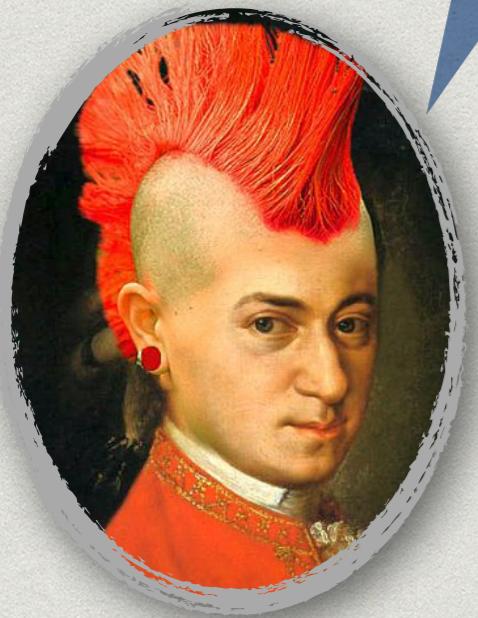
Search w

$UT_1(w)$

$UT_2(w)$

....

$UT_c(w)$



Add (ind_1, \dots, ind_c) to w

Search w

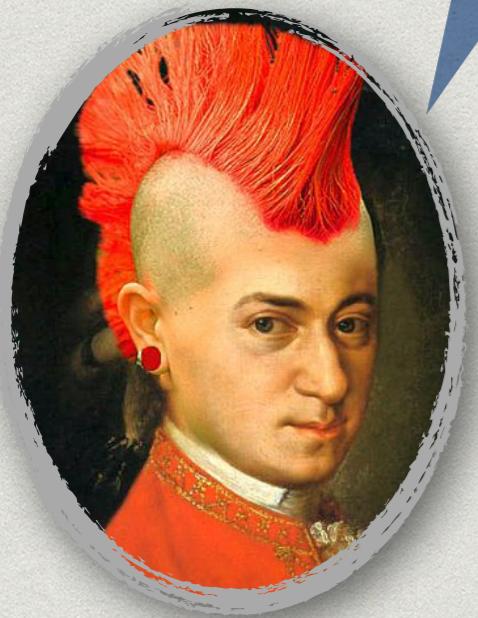
$UT_1(w)$

$UT_2(w)$

...

$UT_c(w)$

$ST(w)$



Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

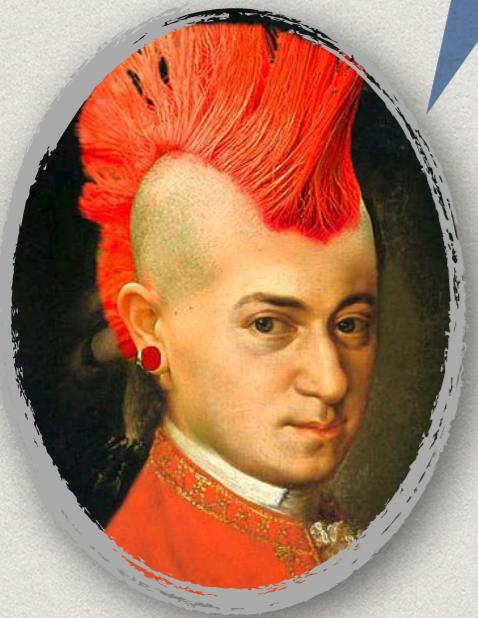
$\text{UT}_1(w)$

$\text{UT}_2(w)$

...

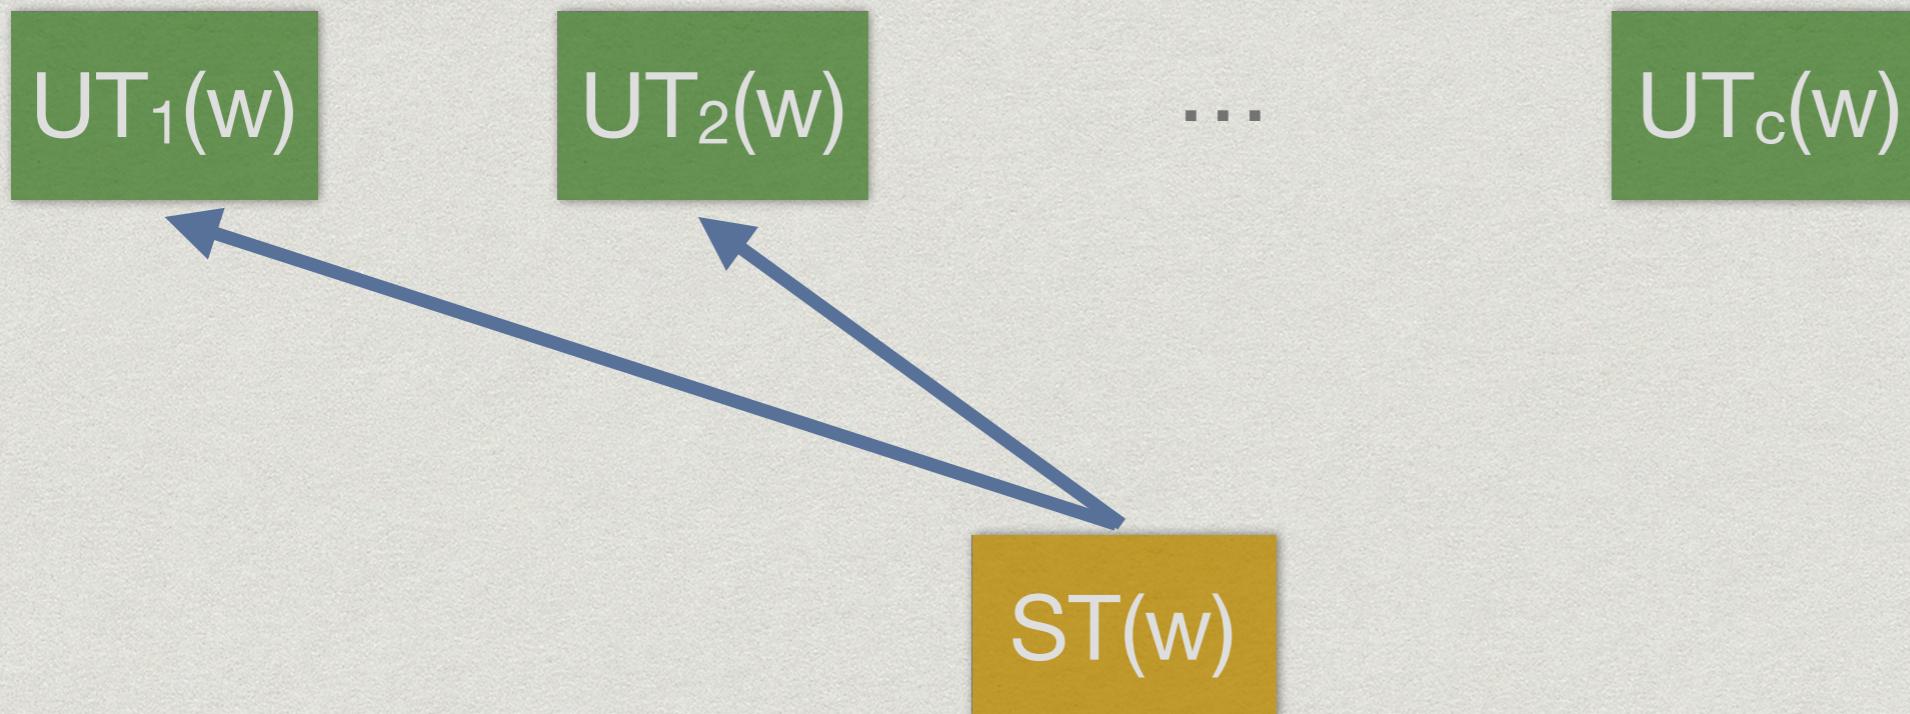
$\text{UT}_c(w)$

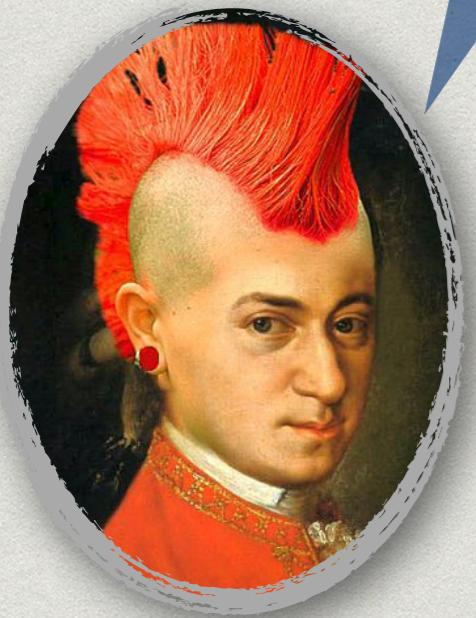
$\text{ST}(w)$



Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

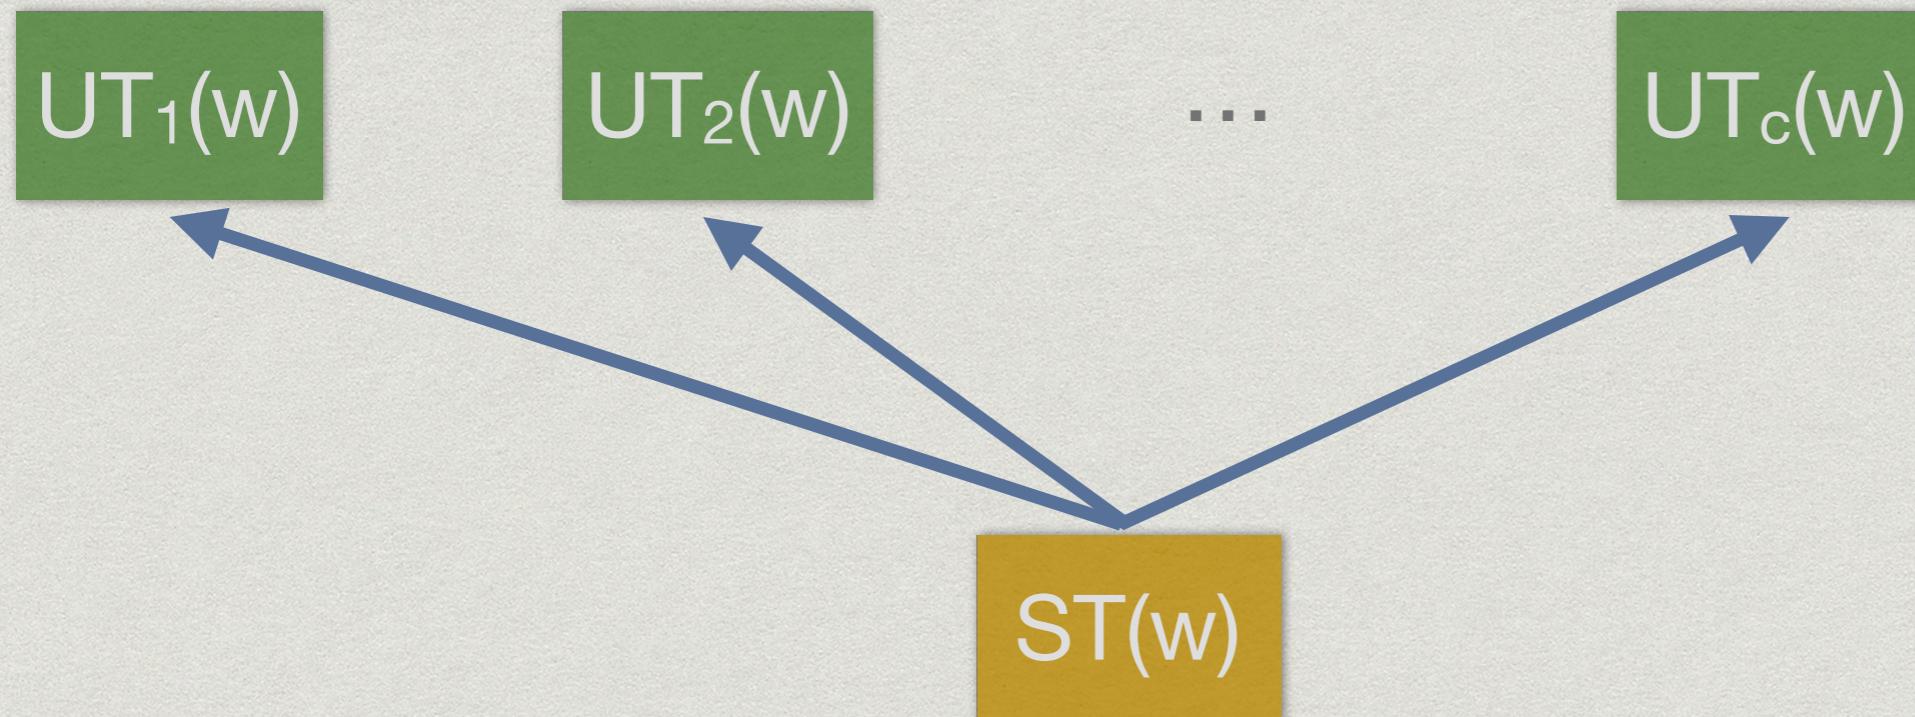
Search w

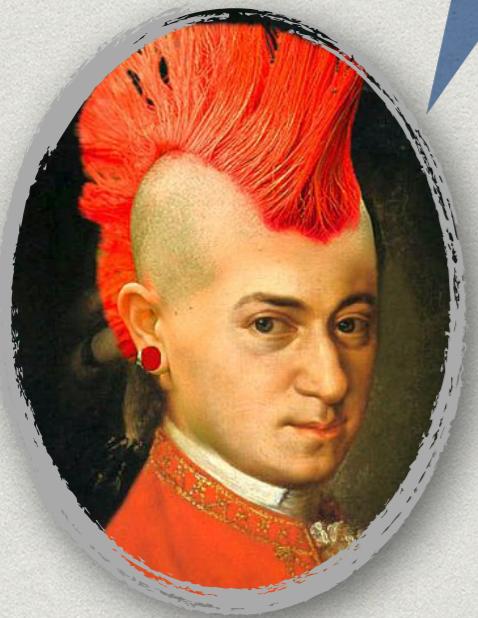




Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

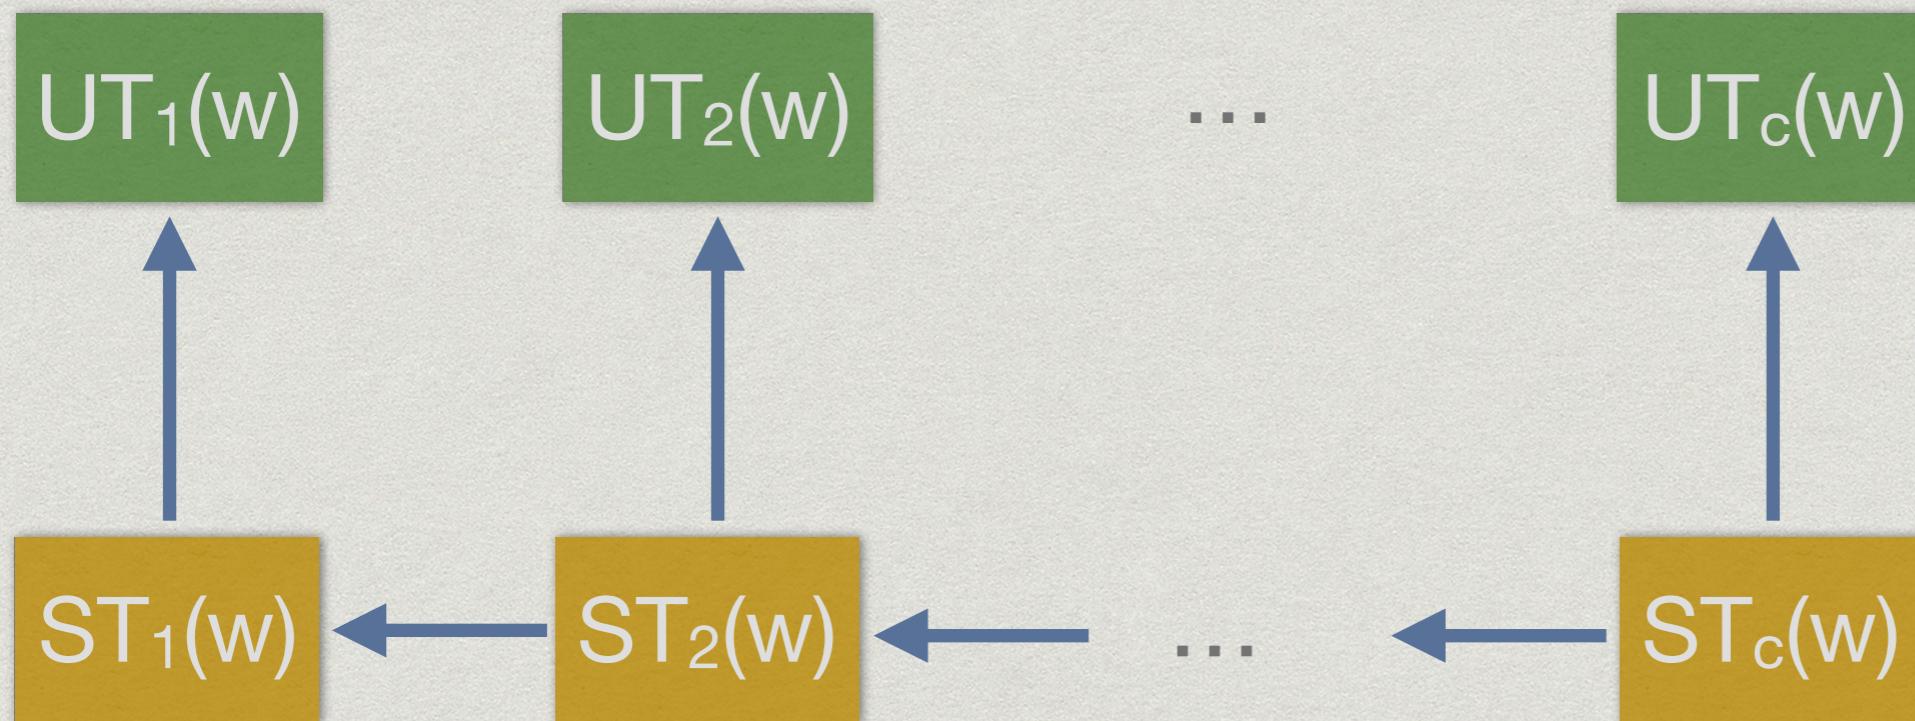
Search w





Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

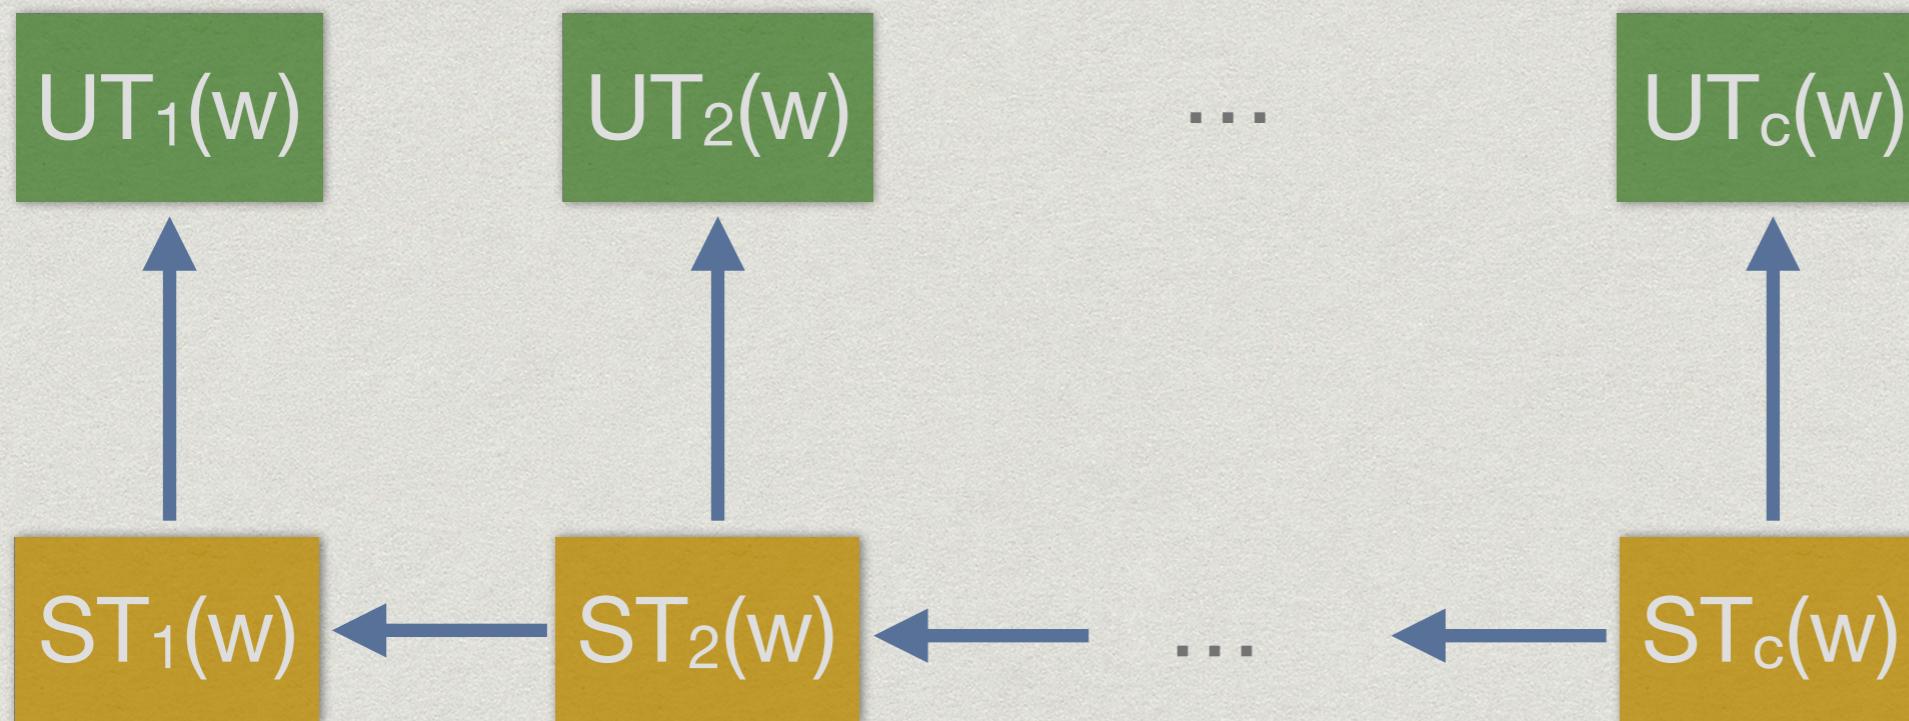


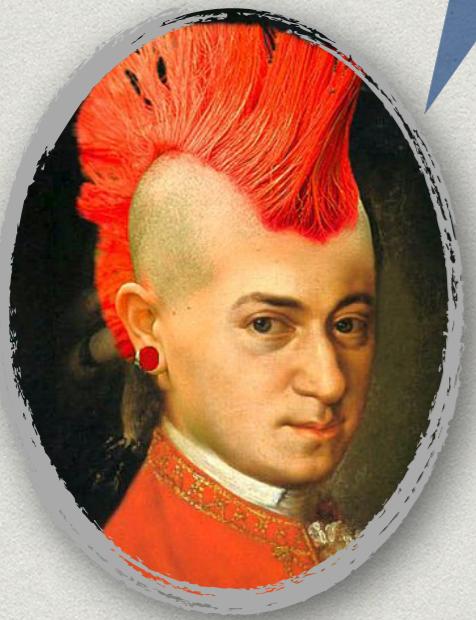


Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

Add ind_{c+1} to w

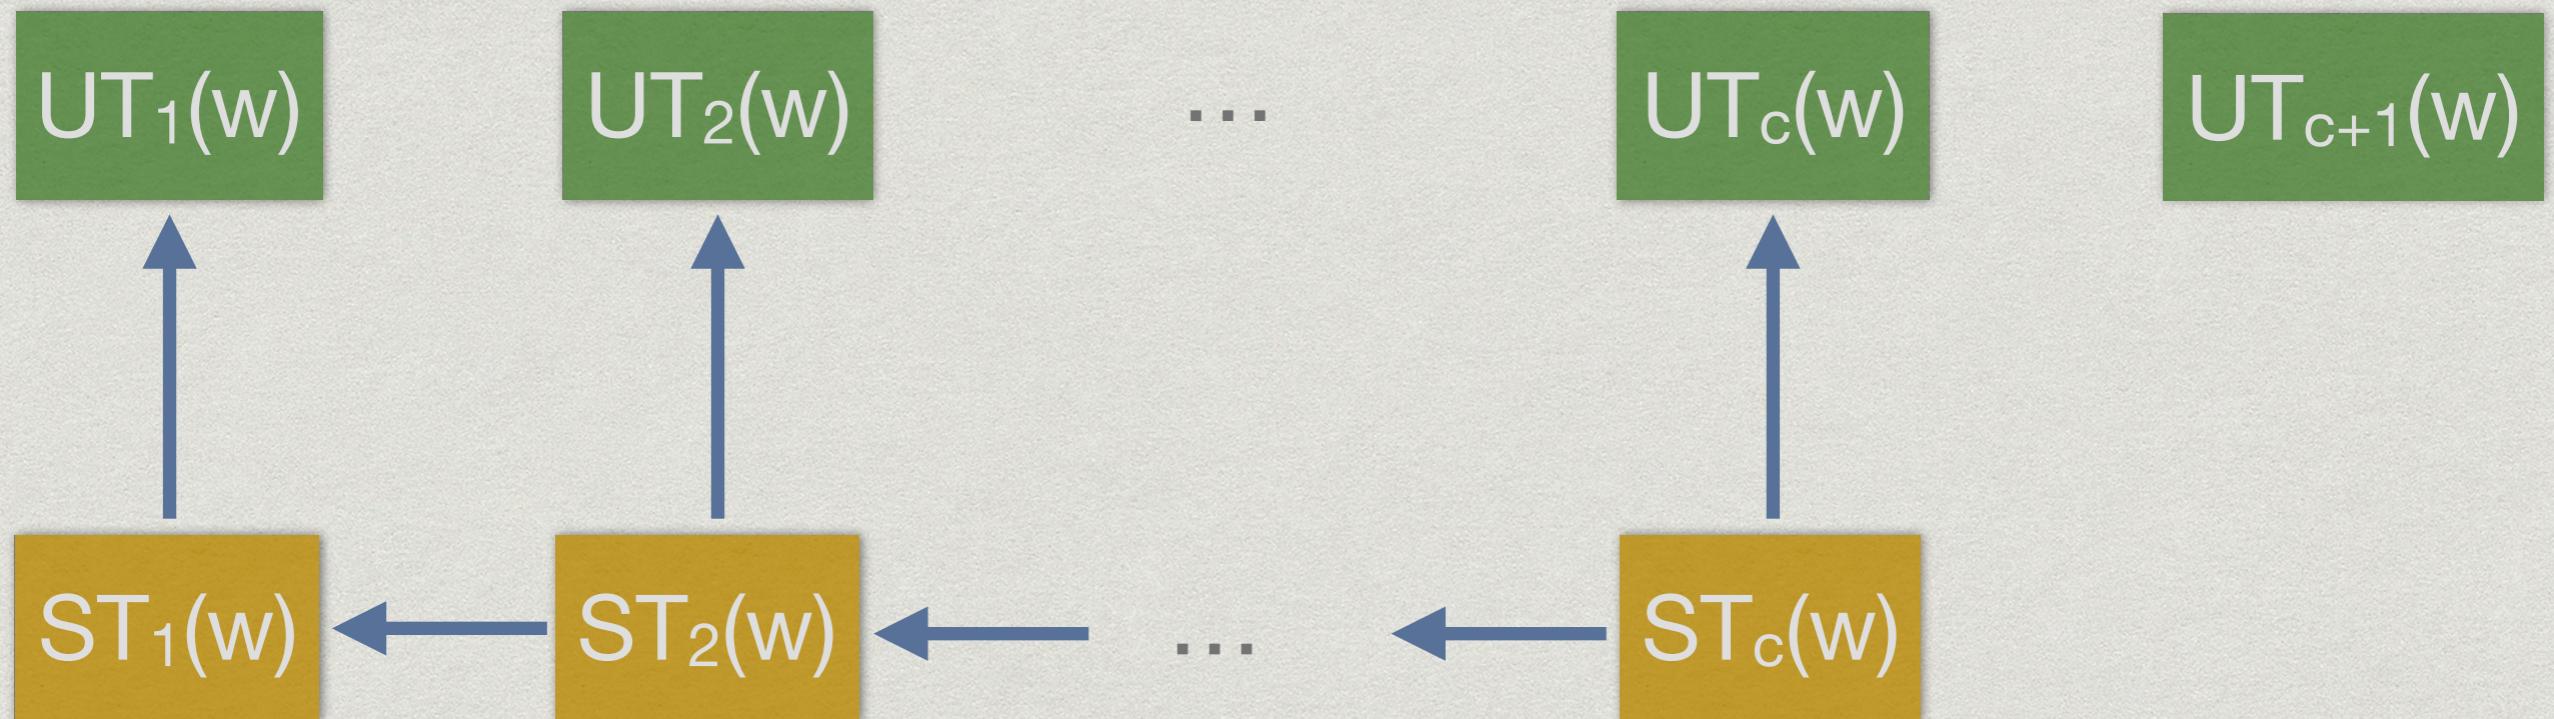


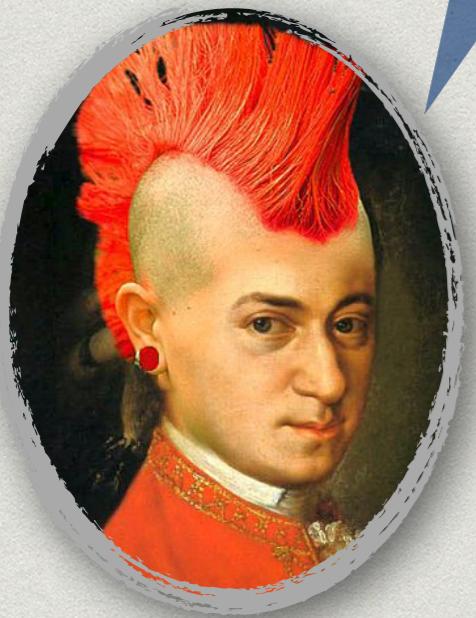


Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

Add ind_{c+1} to w

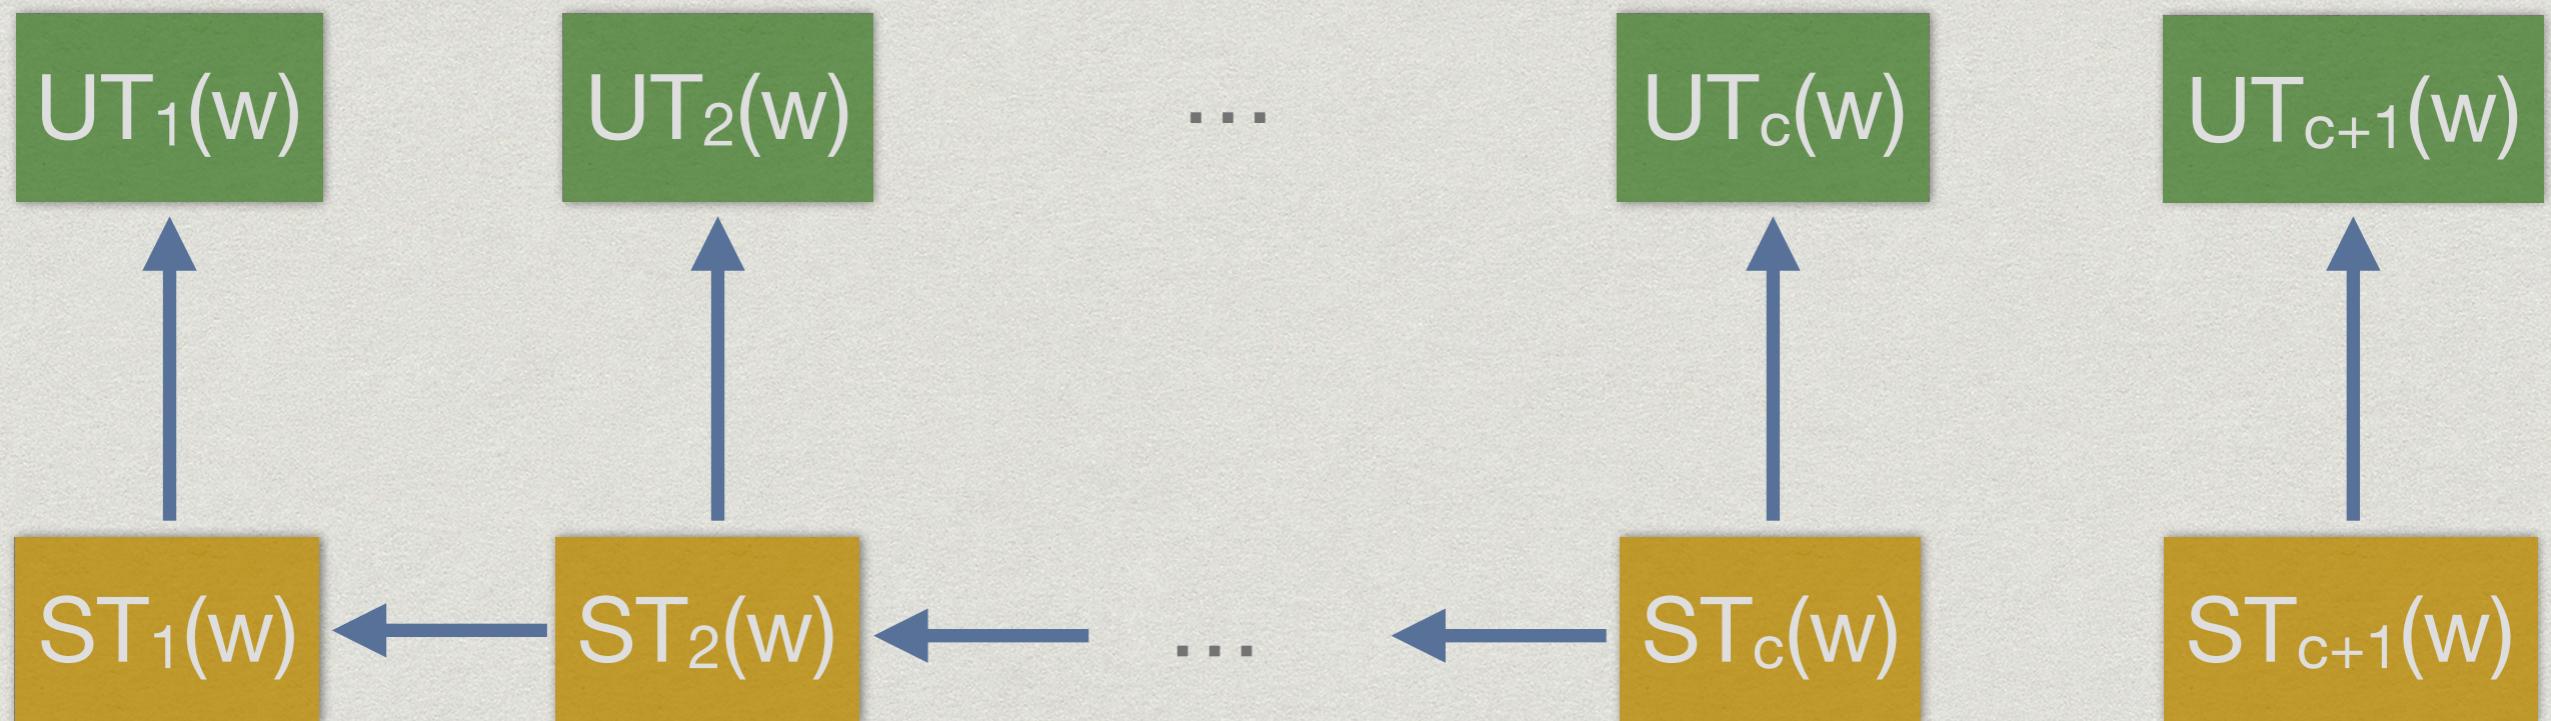


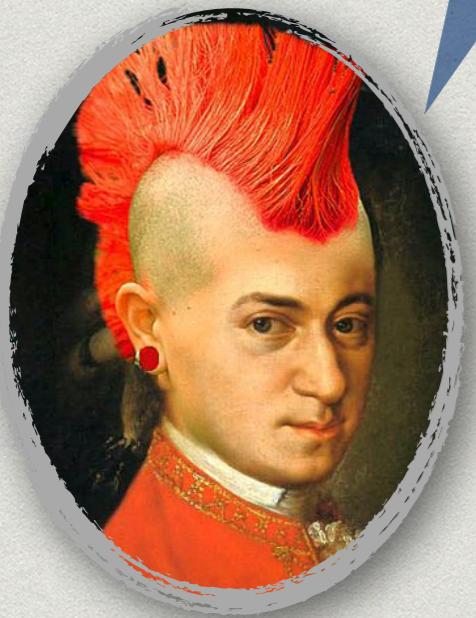


Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

Add ind_{c+1} to w

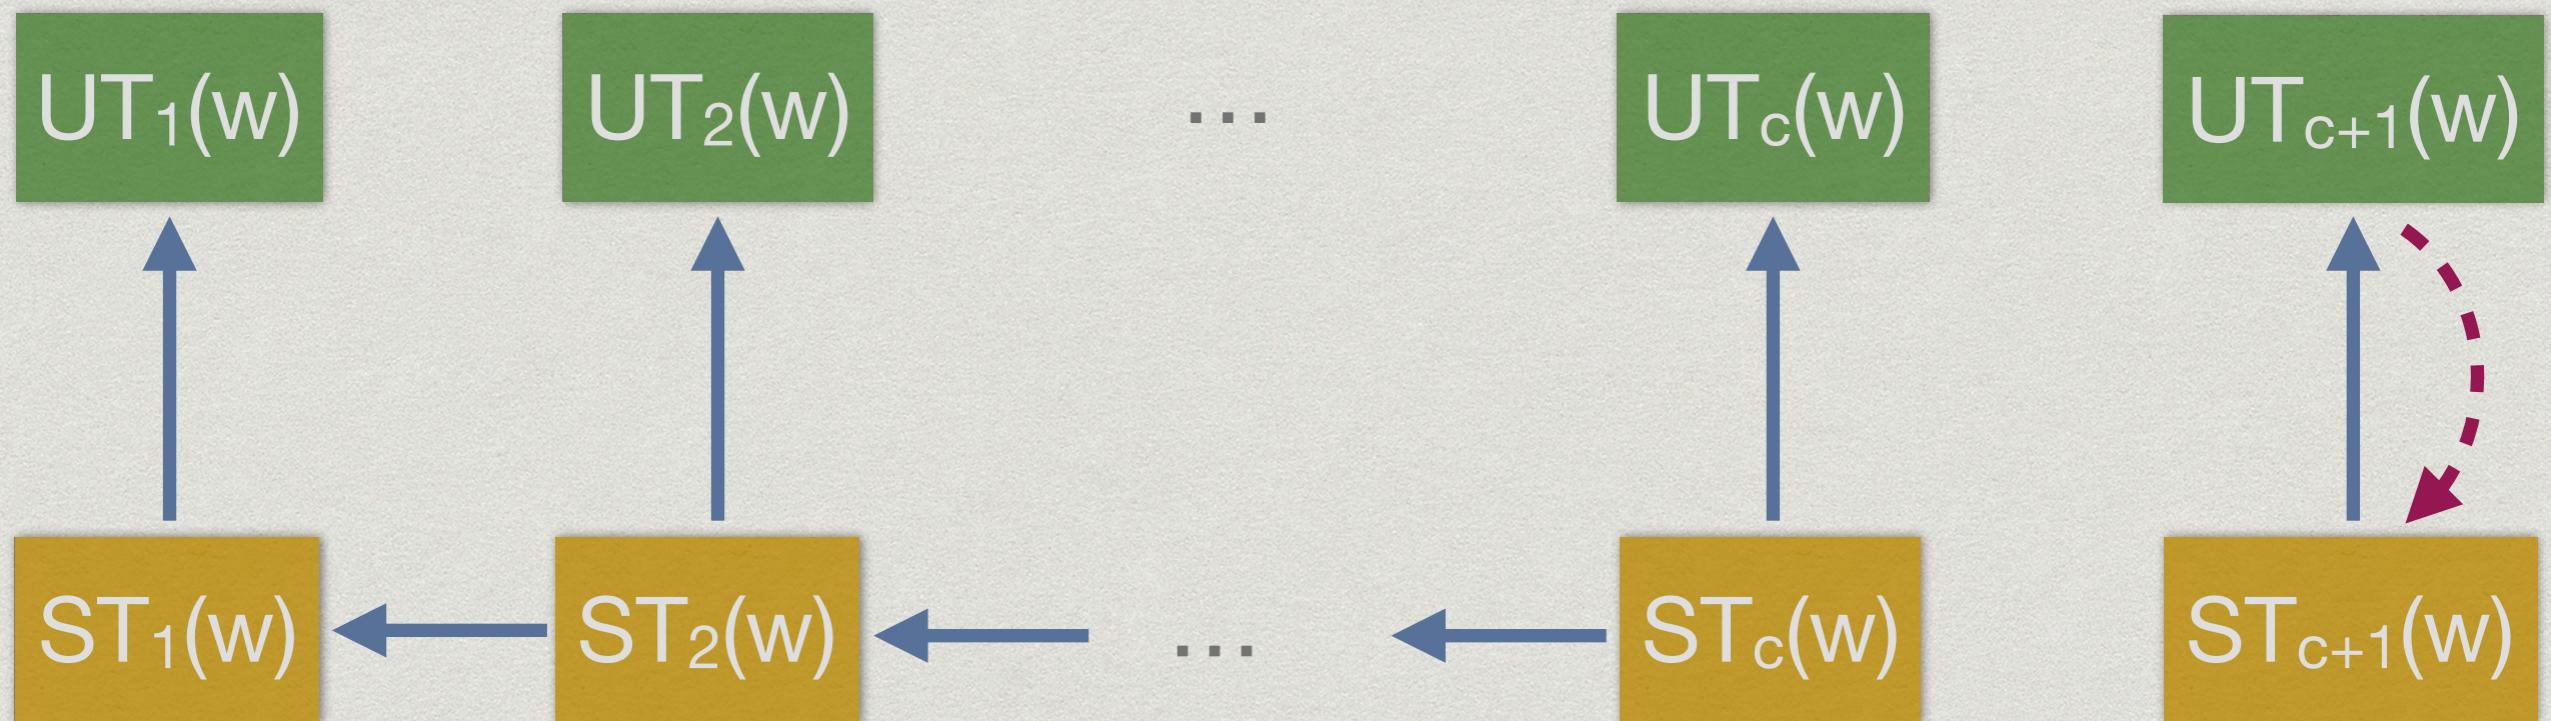


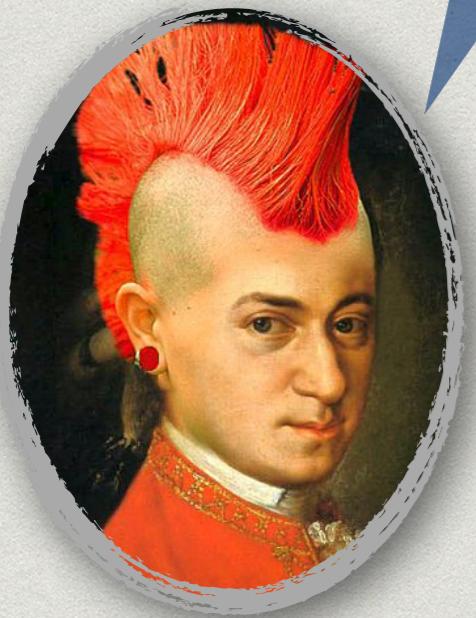


Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

Add ind_{c+1} to w

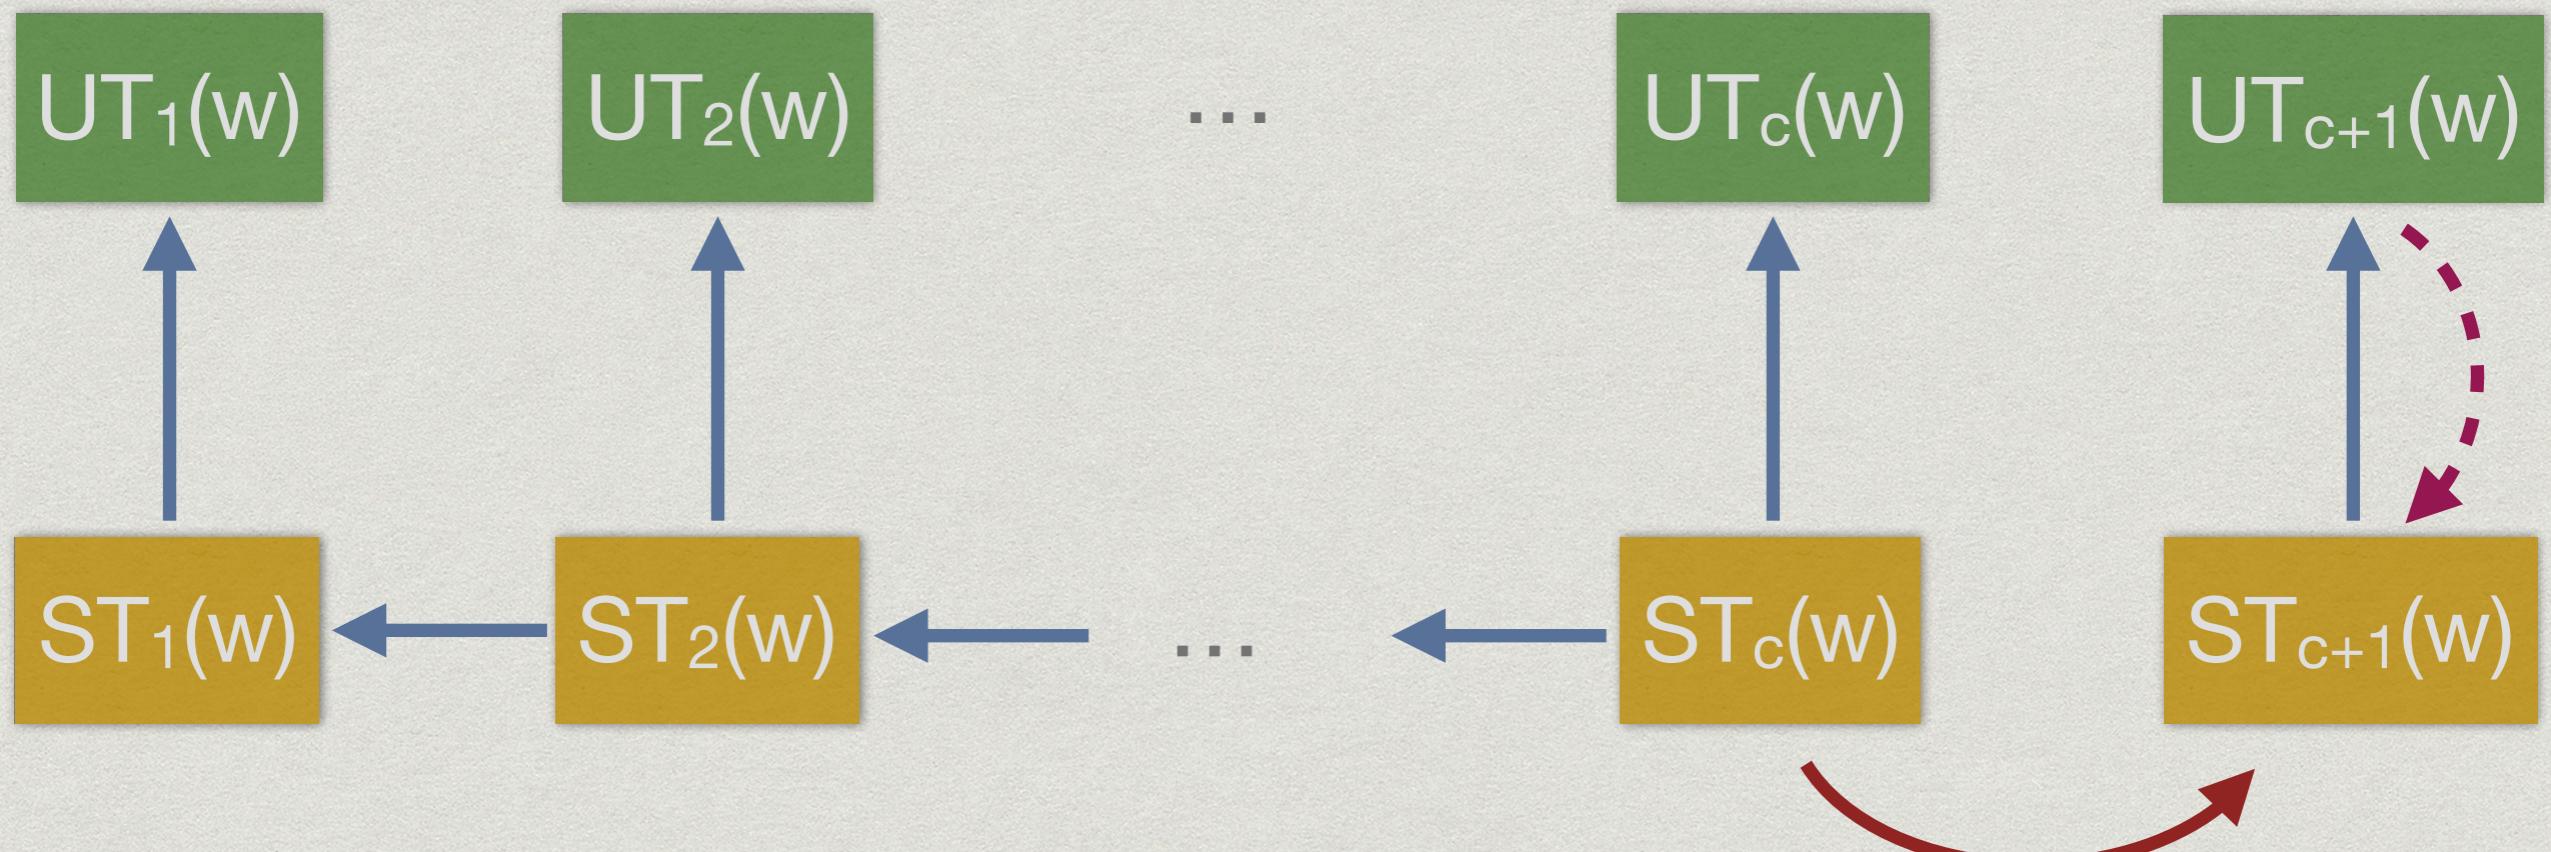


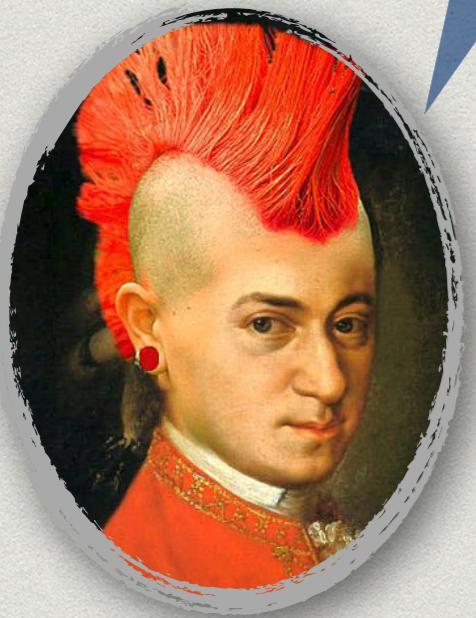


Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

Search w

Add ind_{c+1} to w

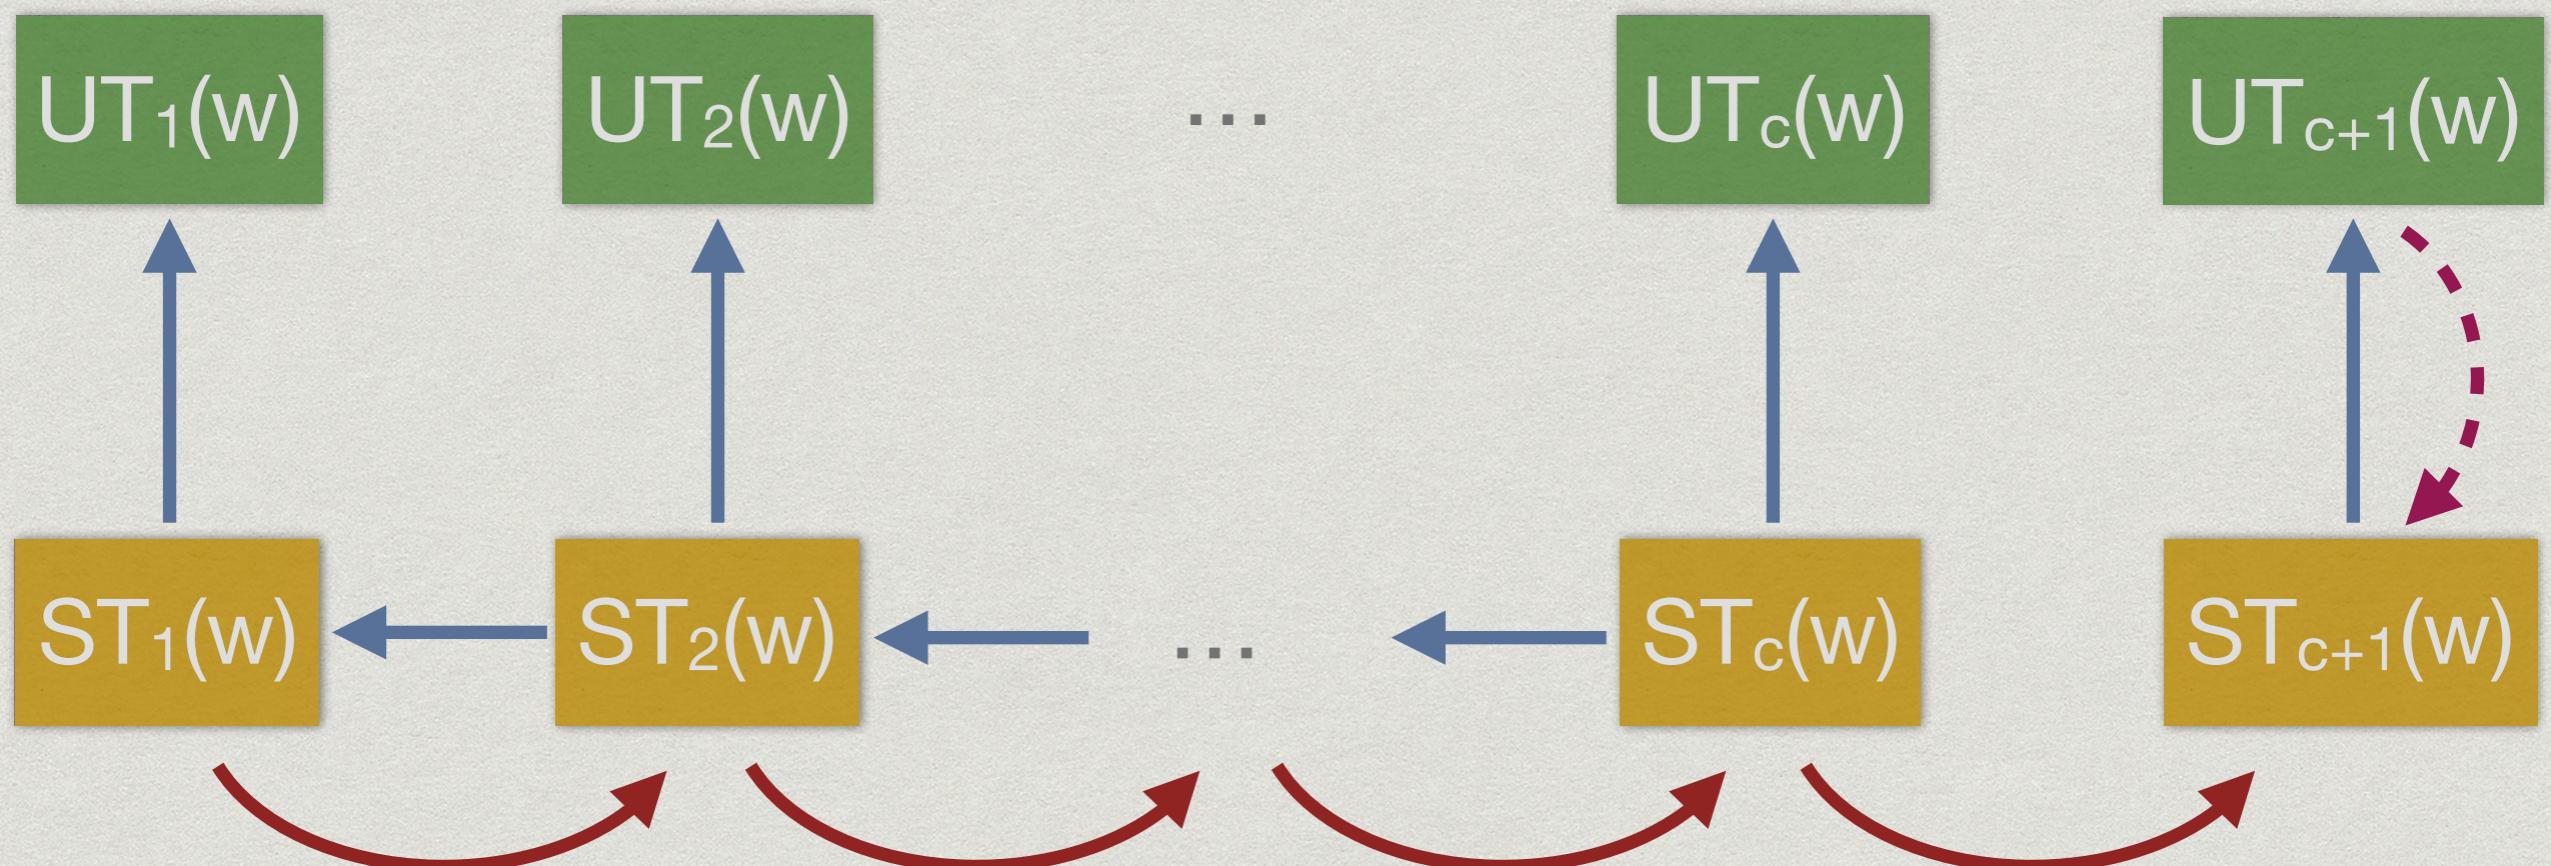


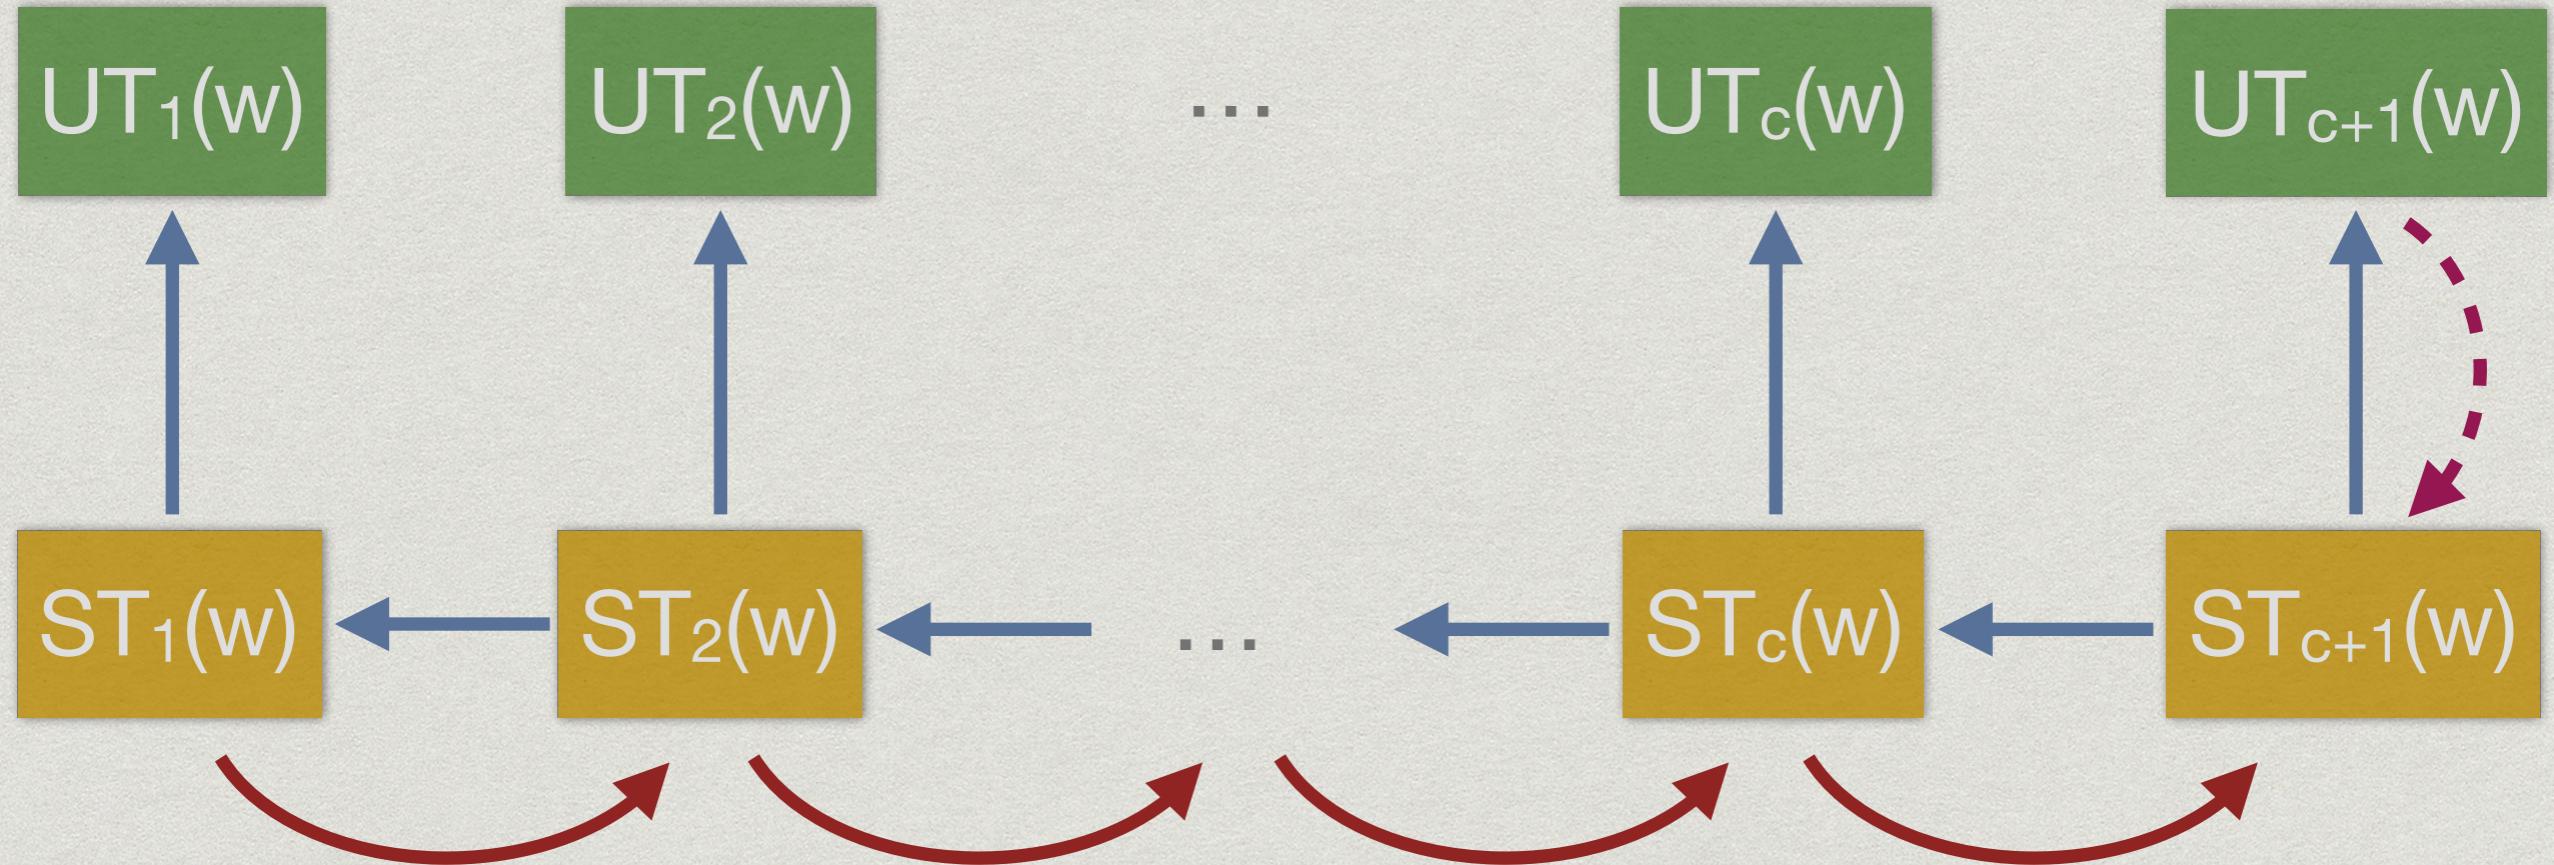


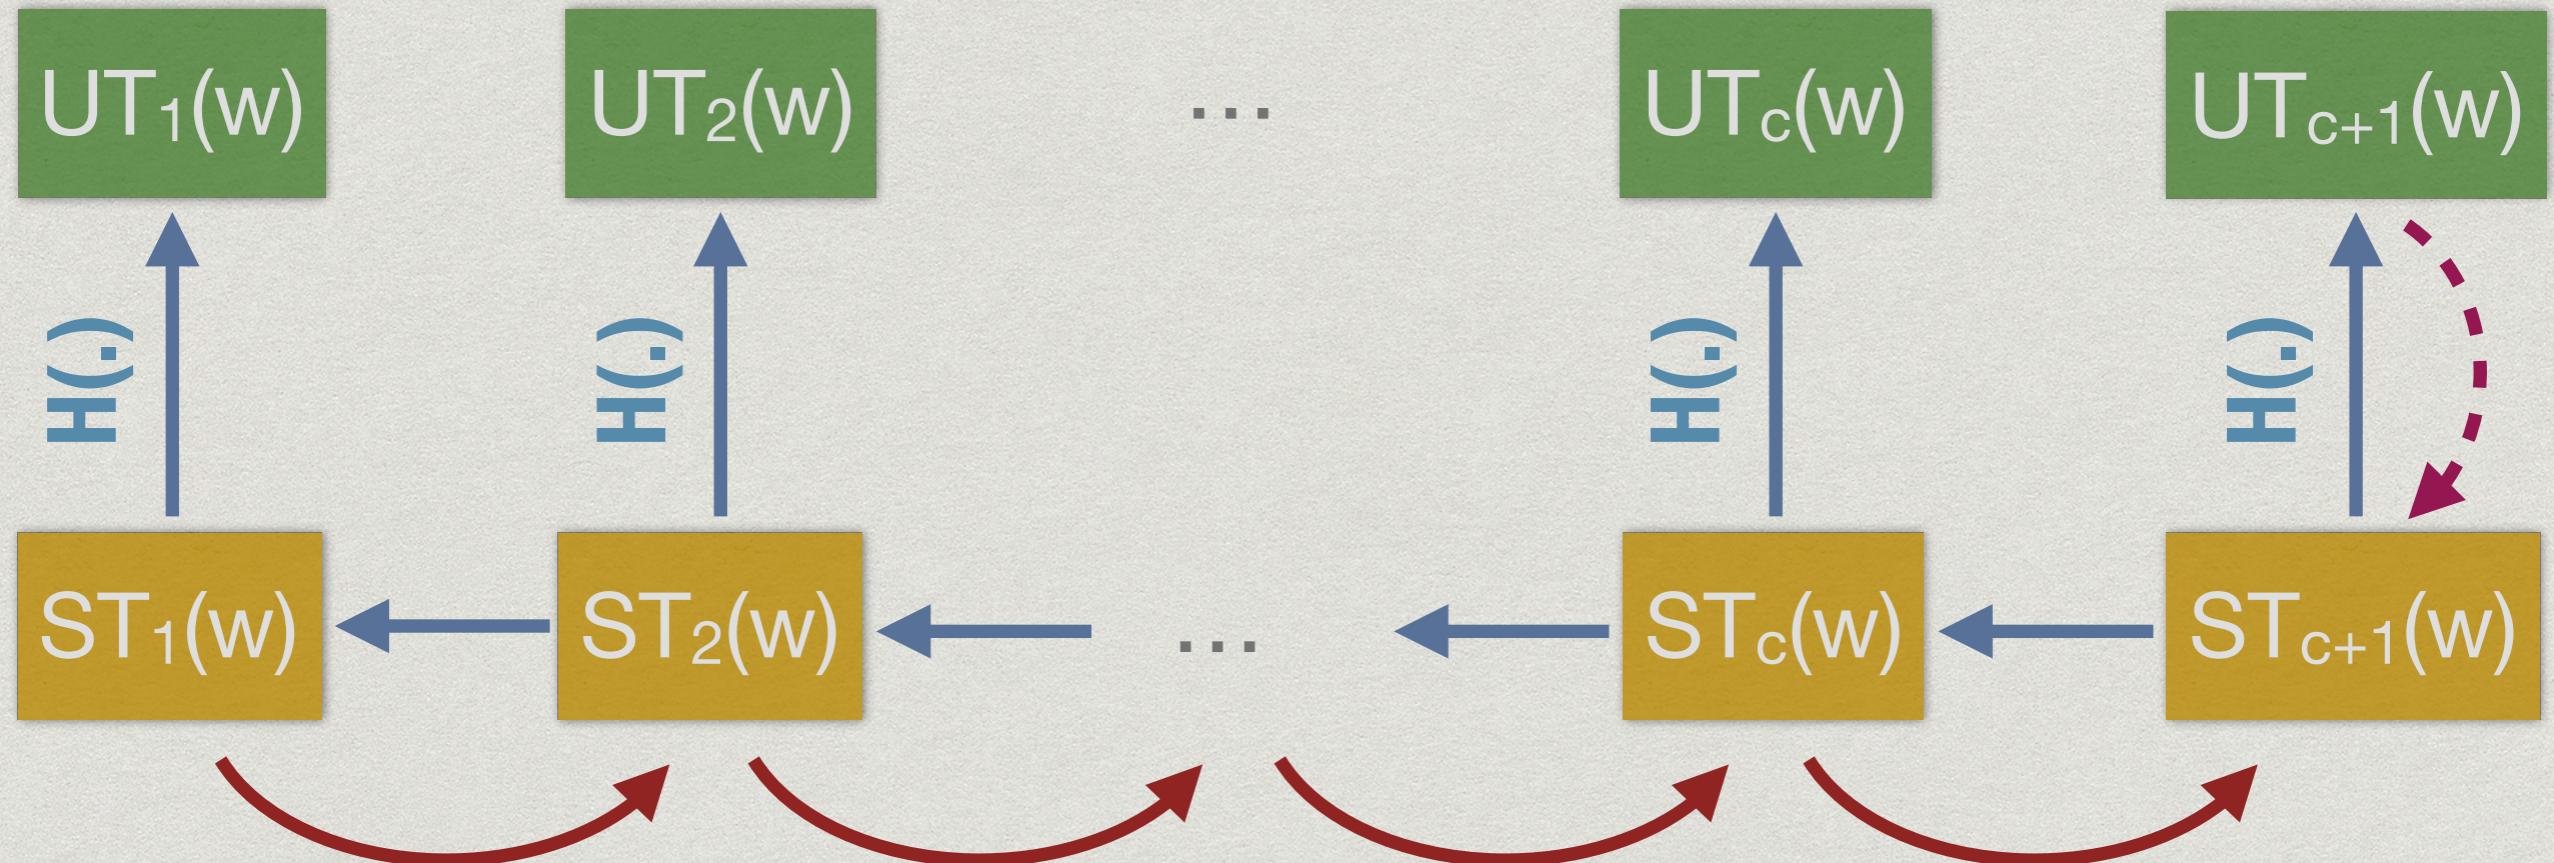
Add $(\text{ind}_1, \dots, \text{ind}_c)$ to w

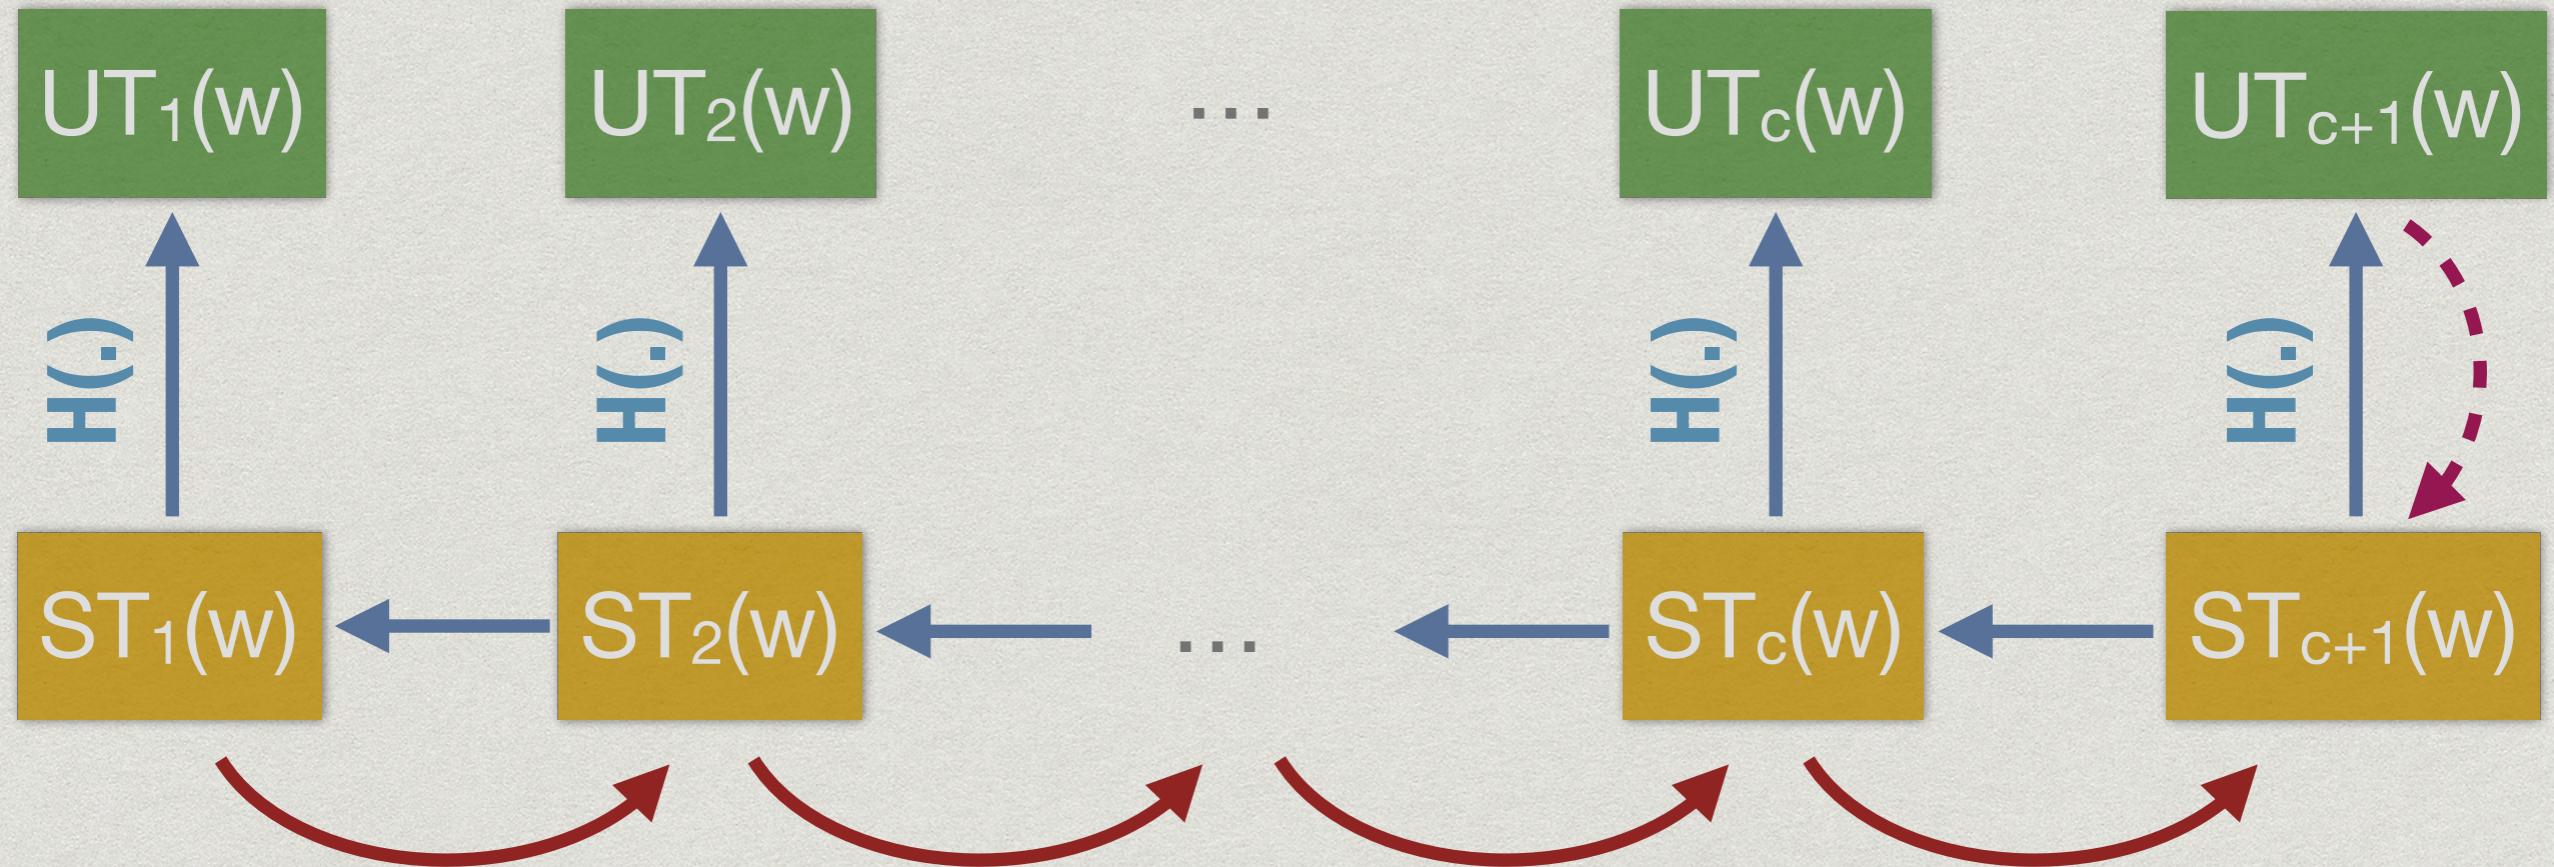
Search w

Add ind_{c+1} to w

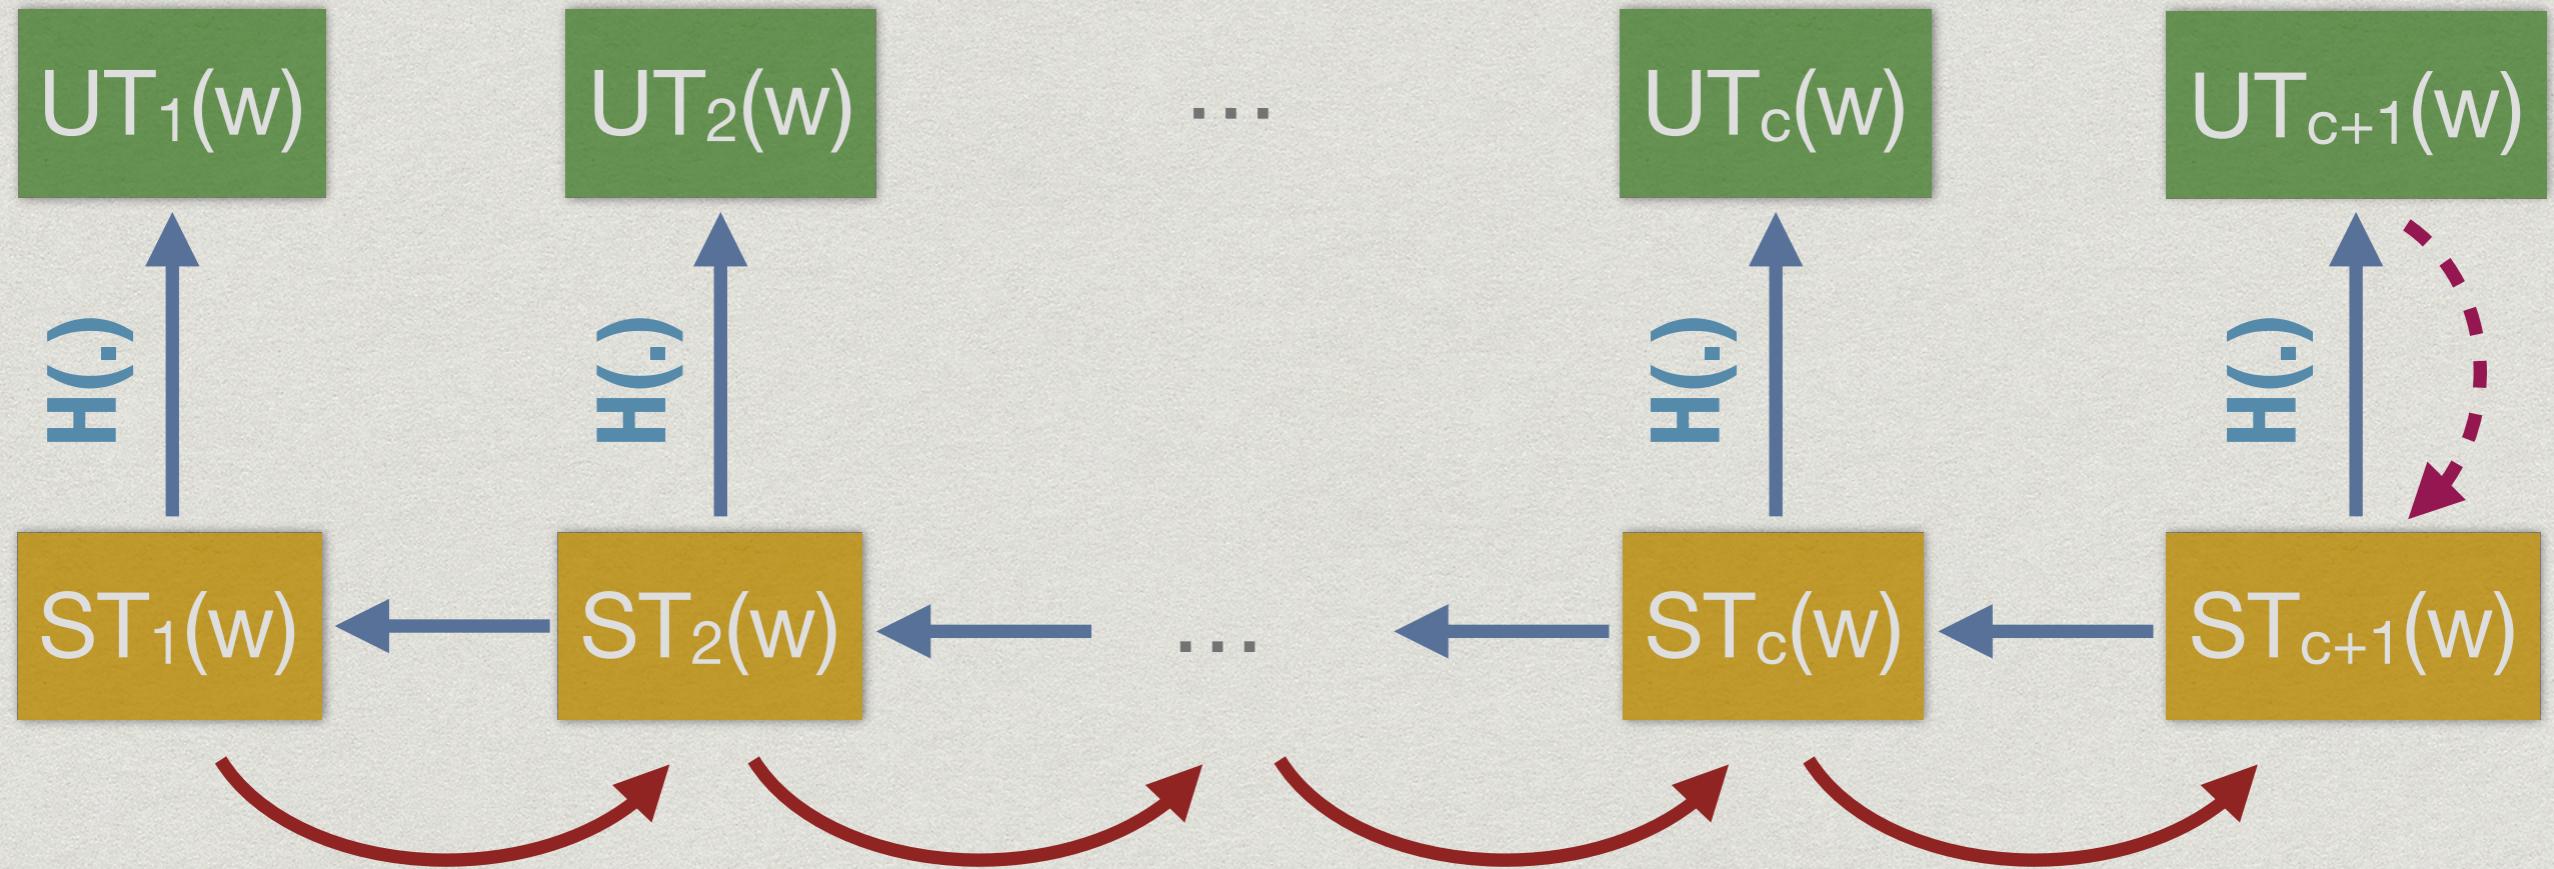




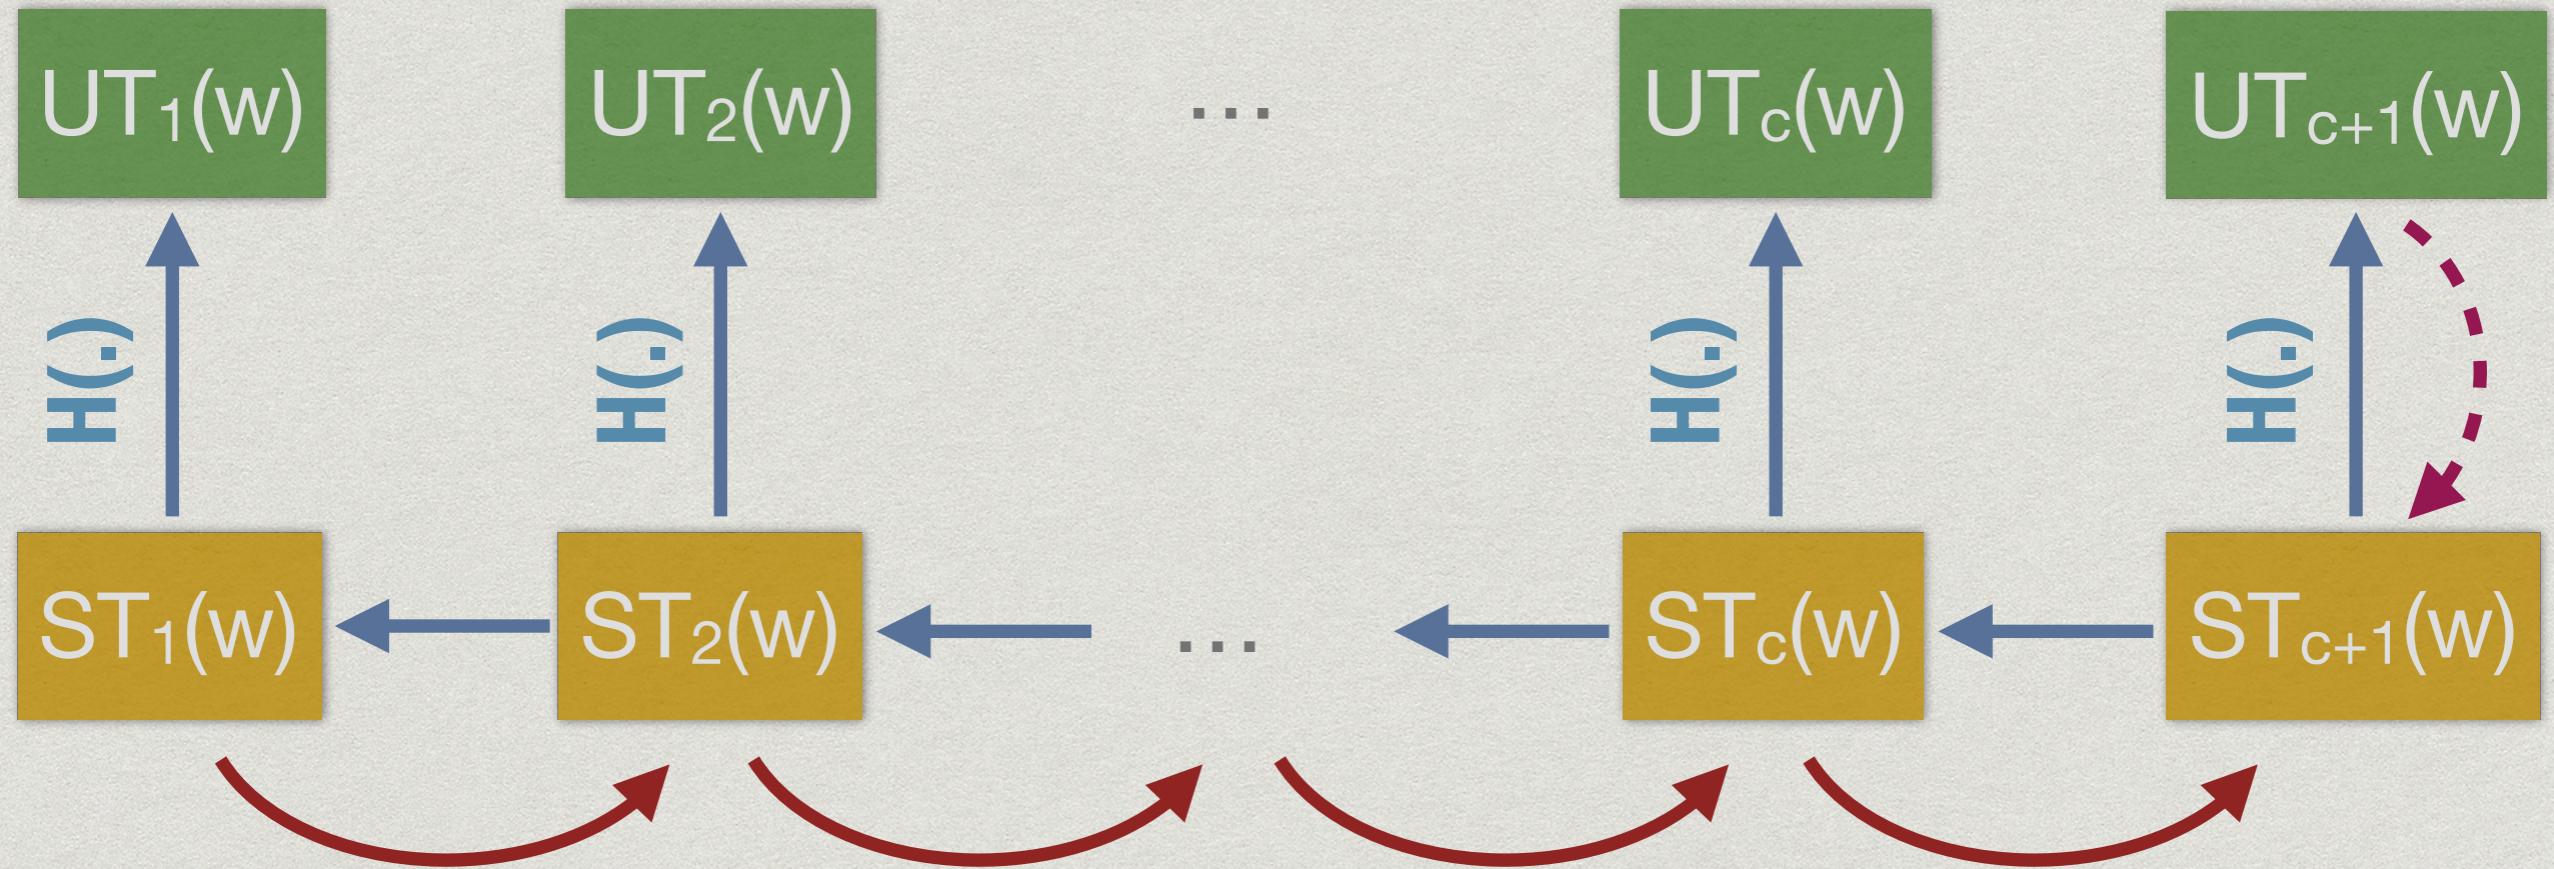




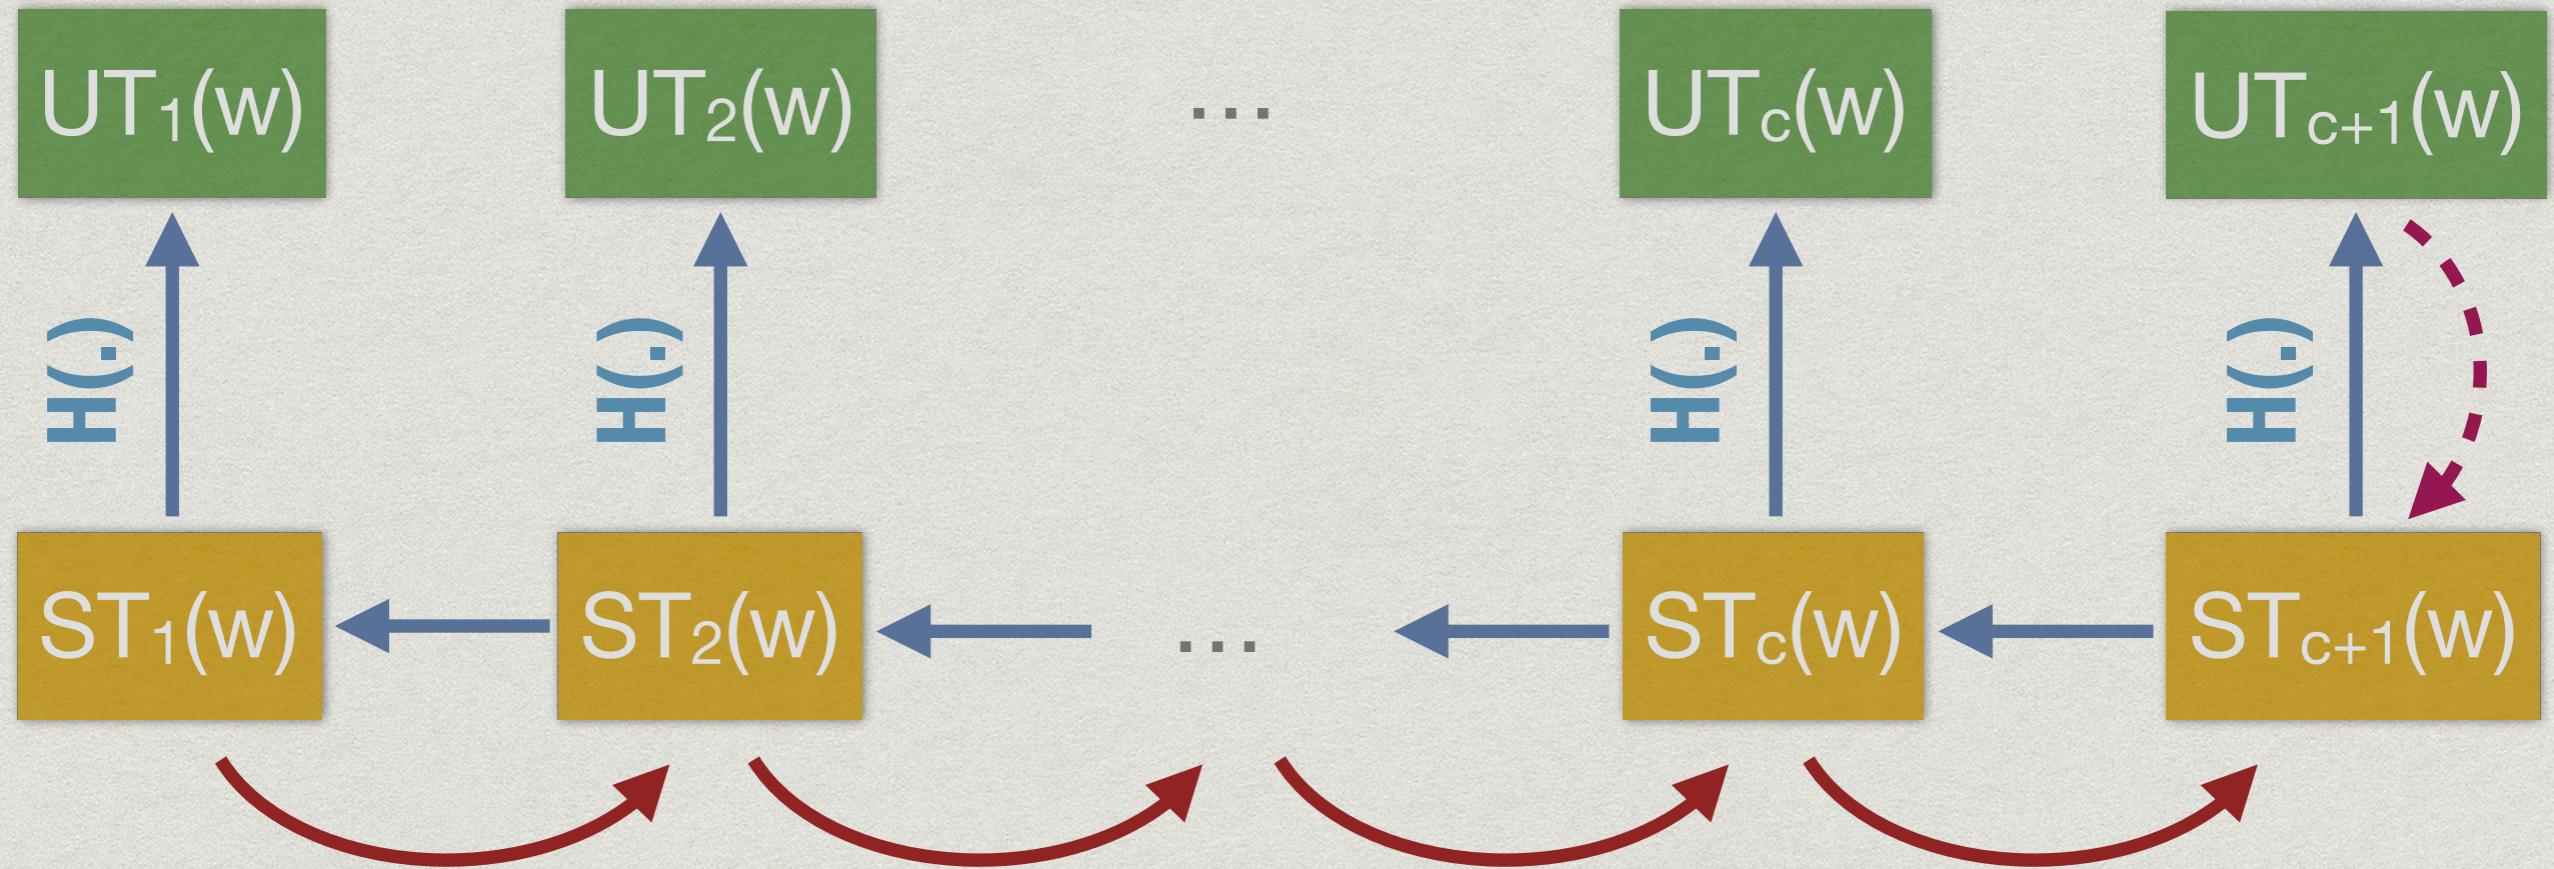
- * Naïve solution: $ST_i(w) = F(K_w, i)$



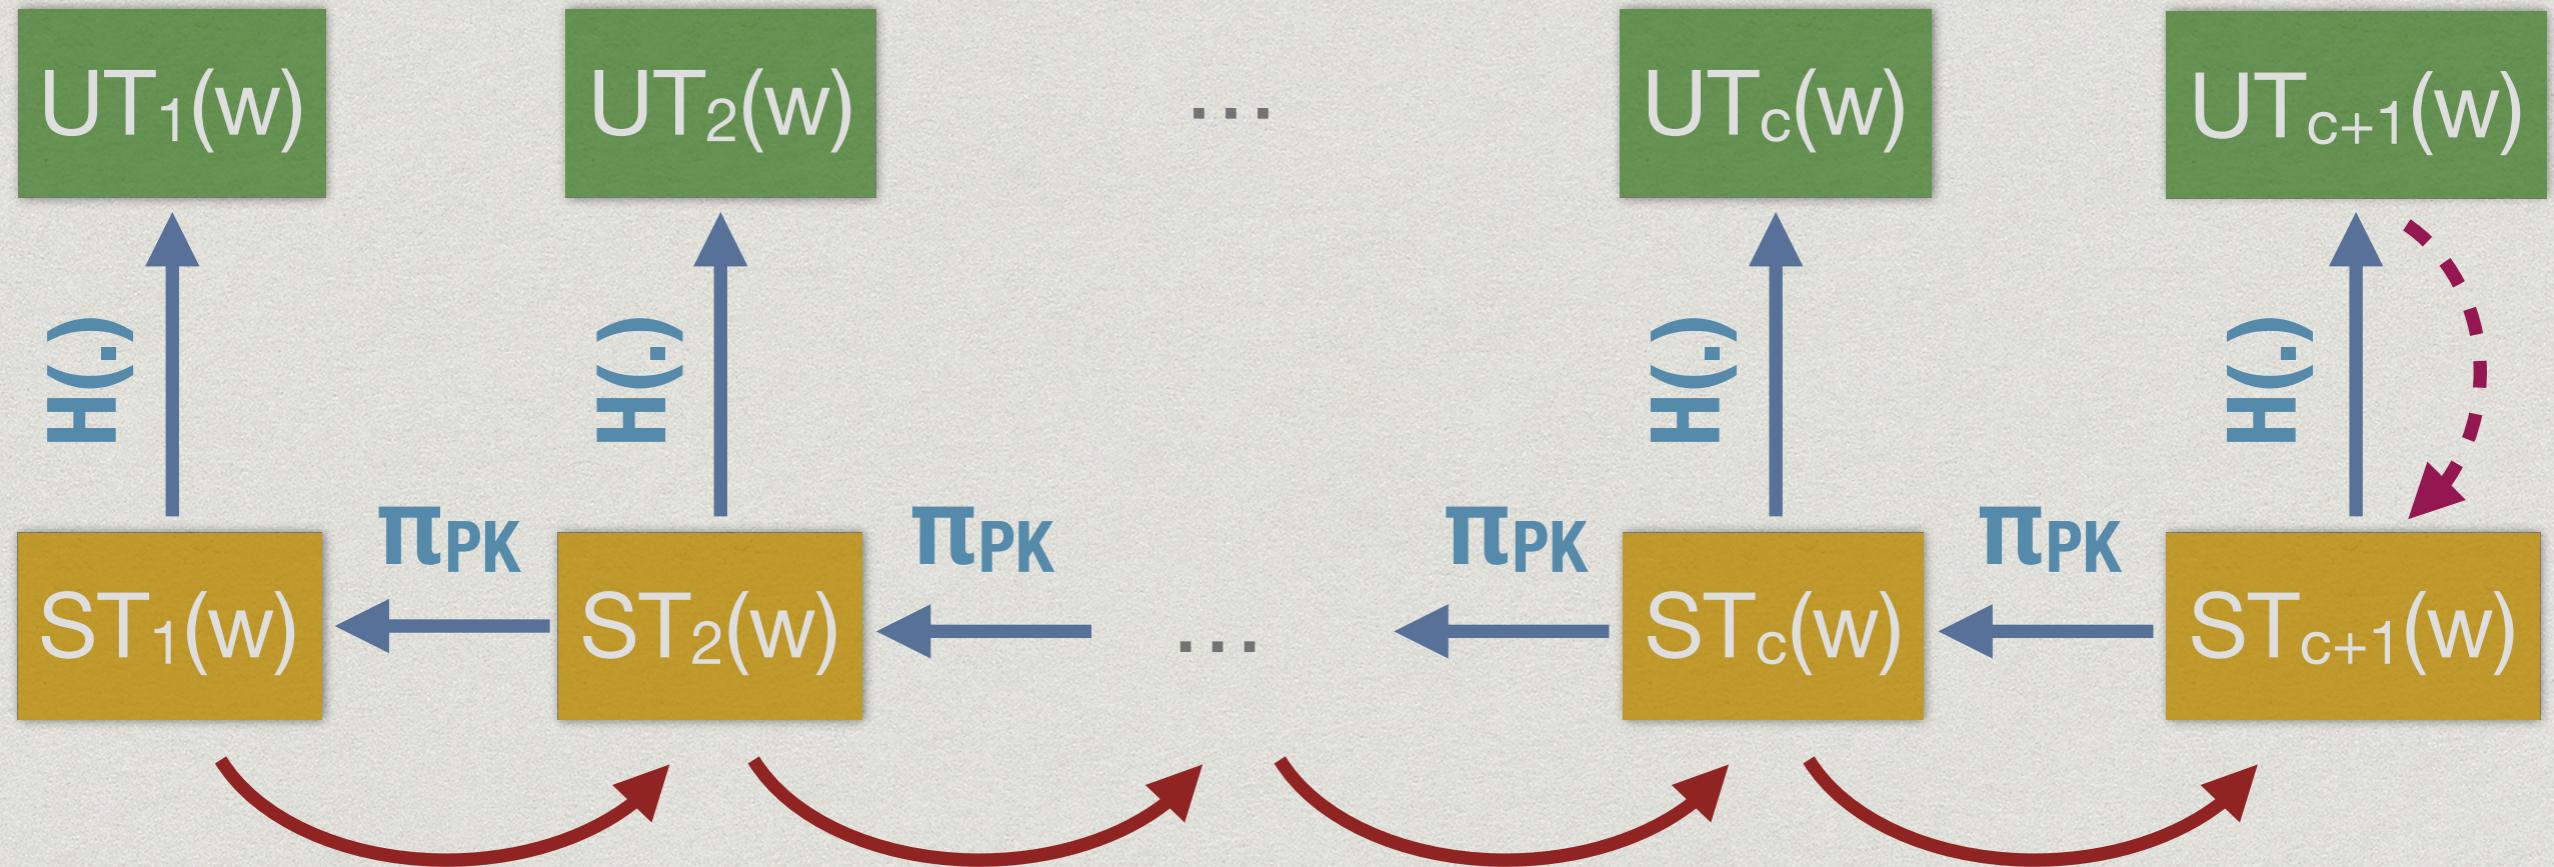
- * Naïve solution: $ST_i(w) = F(K_w, i)$
 - ✗** Sending only K_w is not forward private



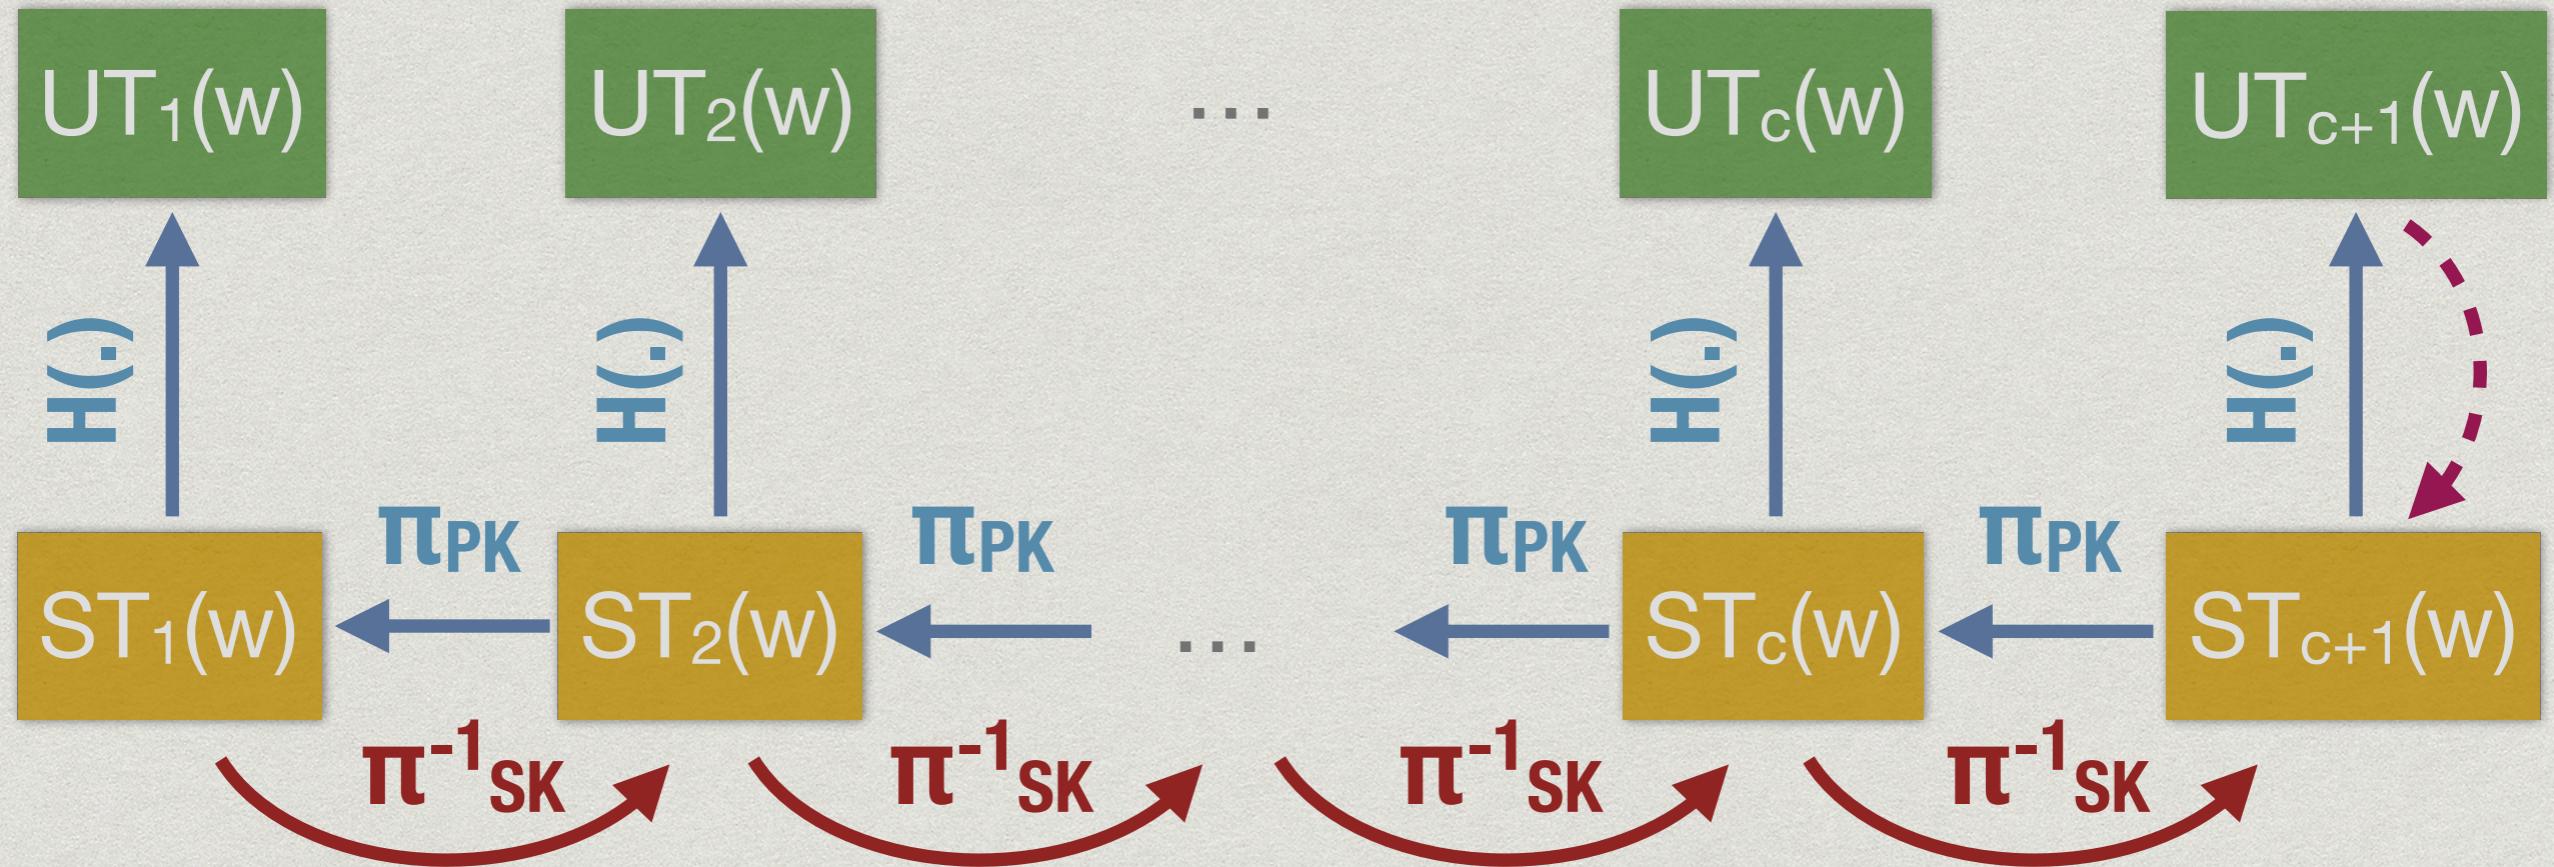
- * Naïve solution: $ST_i(w) = F(K_w, i)$
 - ✗ Sending only K_w is not forward private
 - ✗ O/W client needs to send c tokens



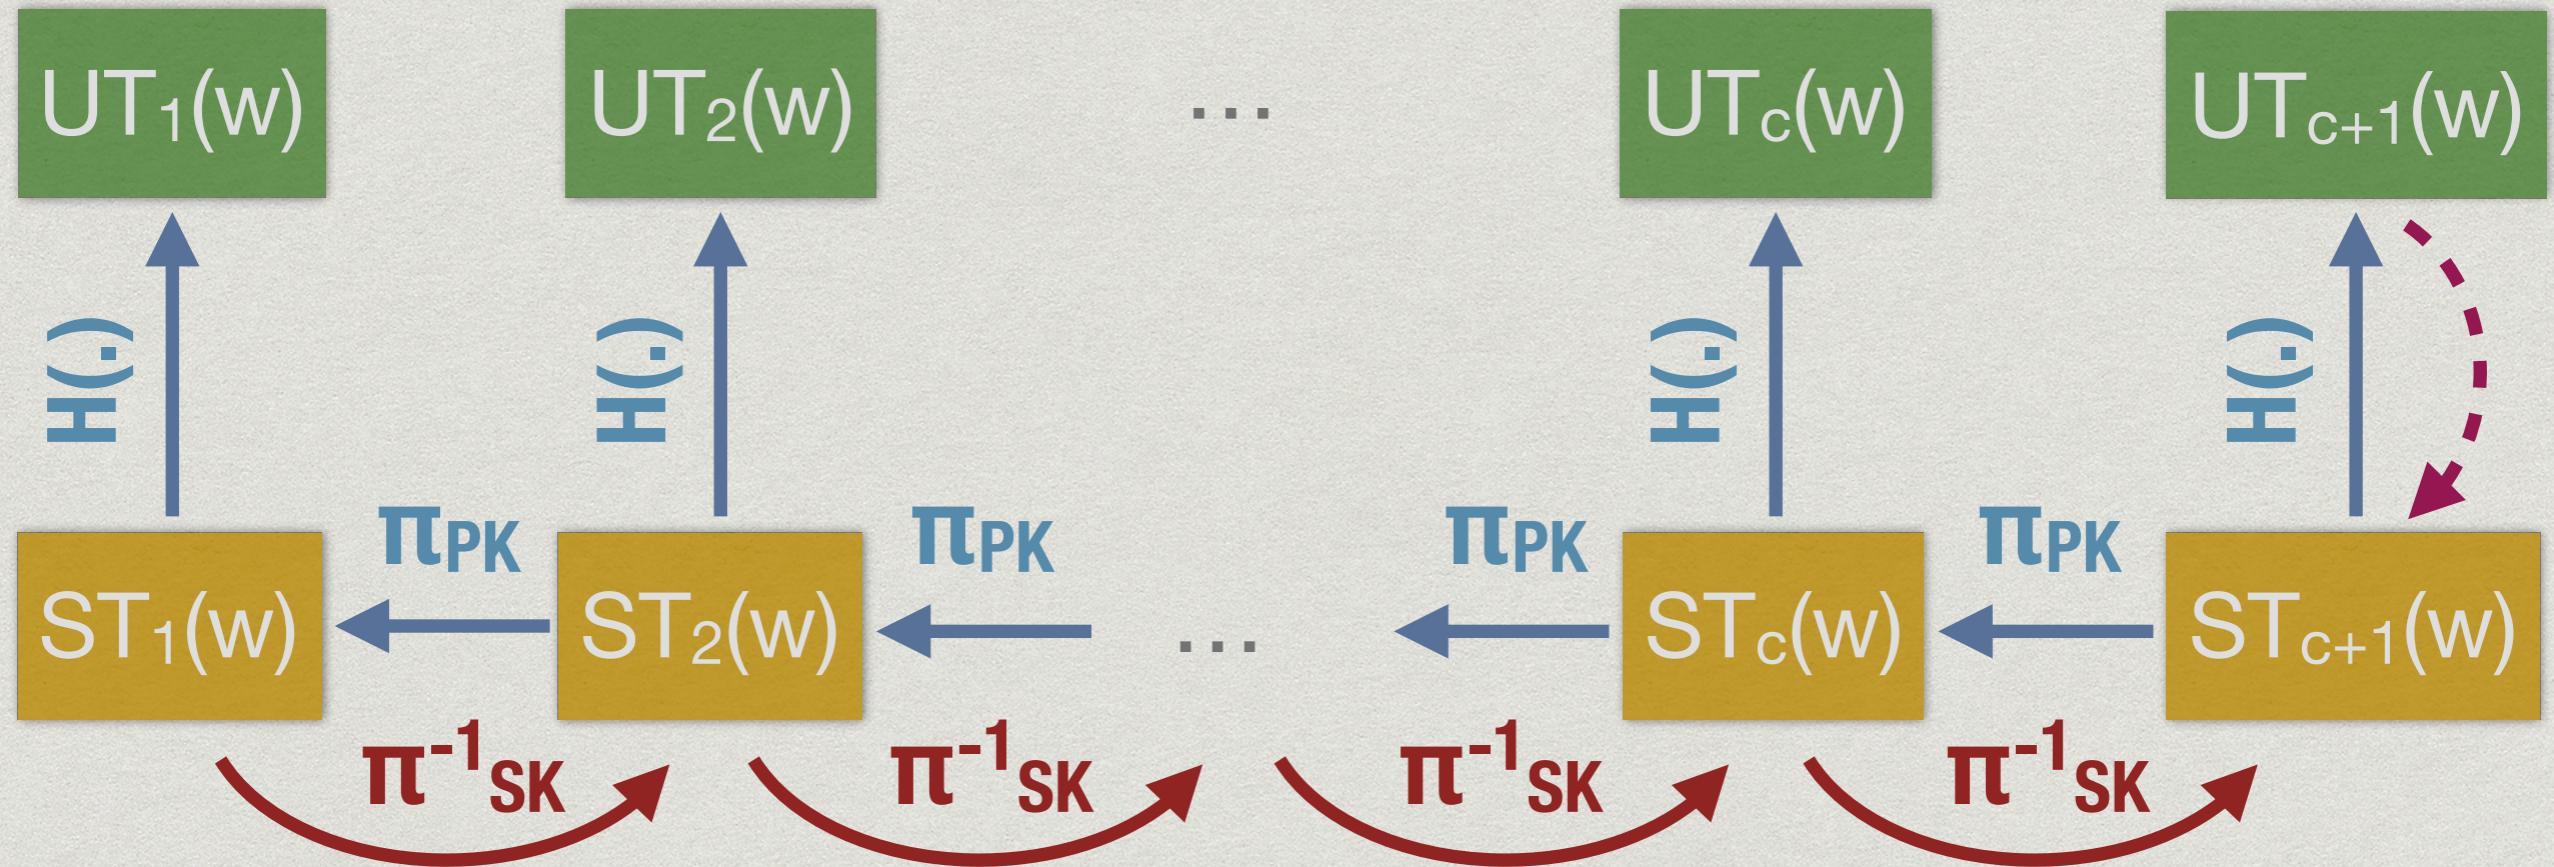
- * Naïve solution: $ST_i(w) = F(K_w, i)$
 - ✗ Sending only K_w is not forward private
 - ✗ O/W client needs to send c tokens
- * Use a trapdoor permutation

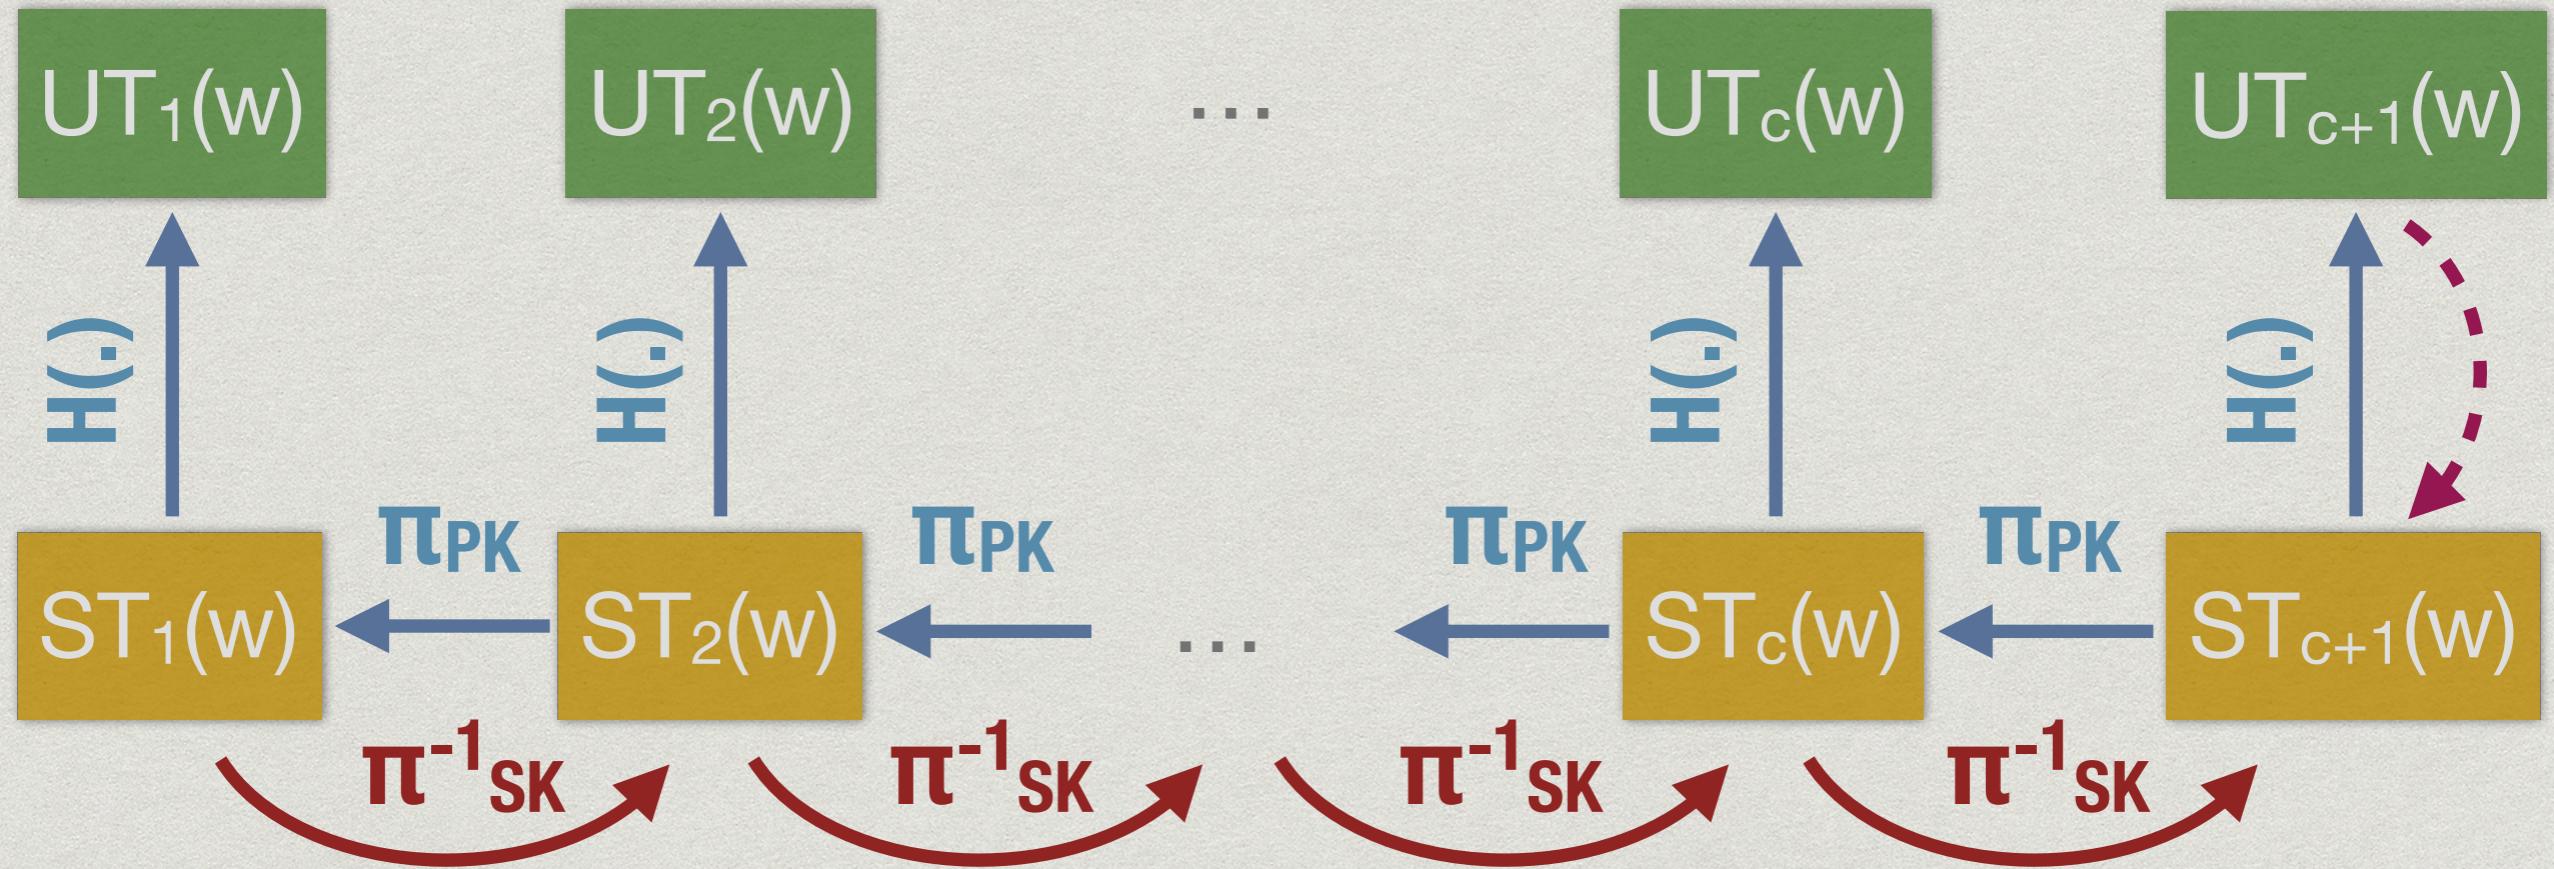


- * Naïve solution: $ST_i(w) = F(K_w, i)$
 - ✗ Sending only K_w is not forward private
 - ✗ O/W client needs to send c tokens
- * Use a trapdoor permutation

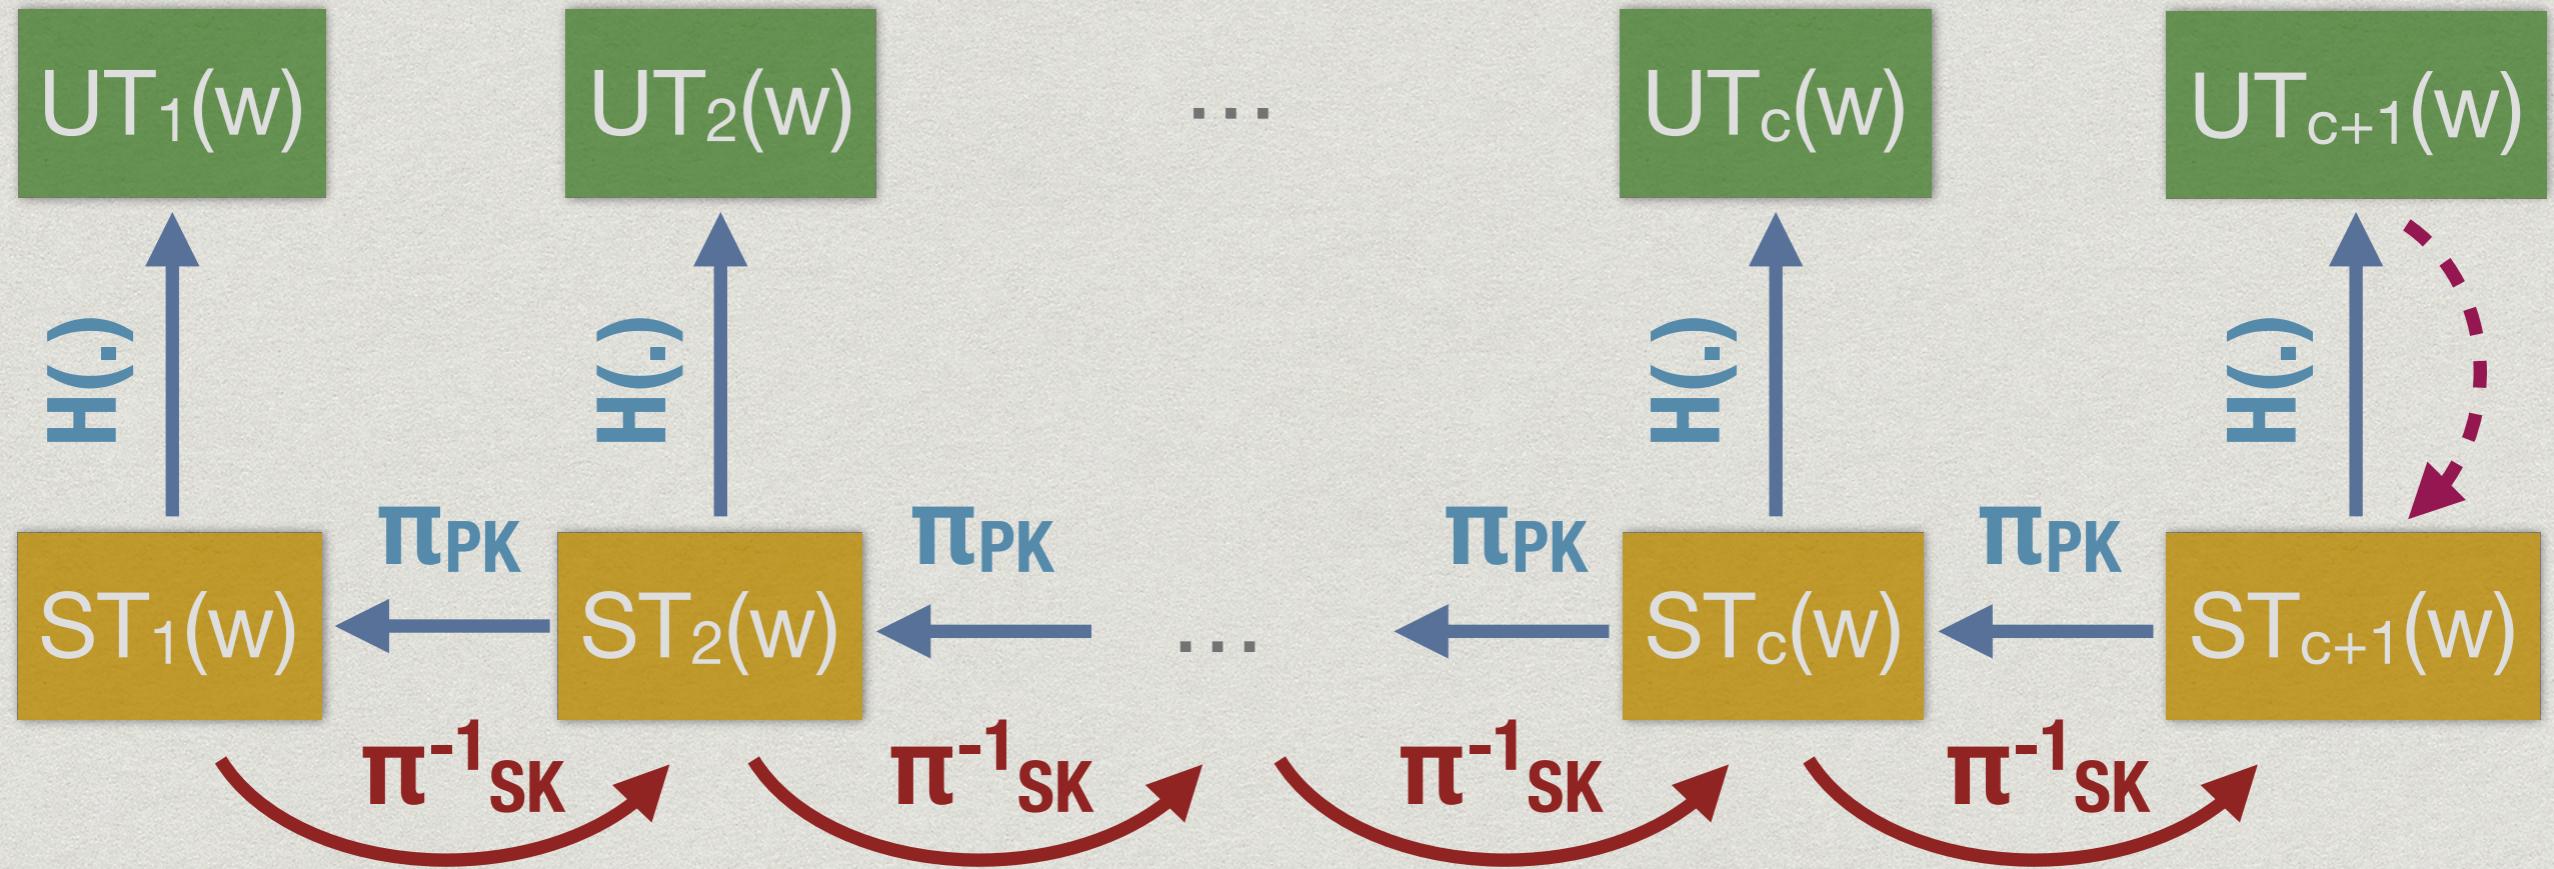


- * Naïve solution: $ST_i(w) = F(K_w, i)$
 - ✗ Sending only K_w is not forward private
 - ✗ O/W client needs to send c tokens
- * Use a trapdoor permutation

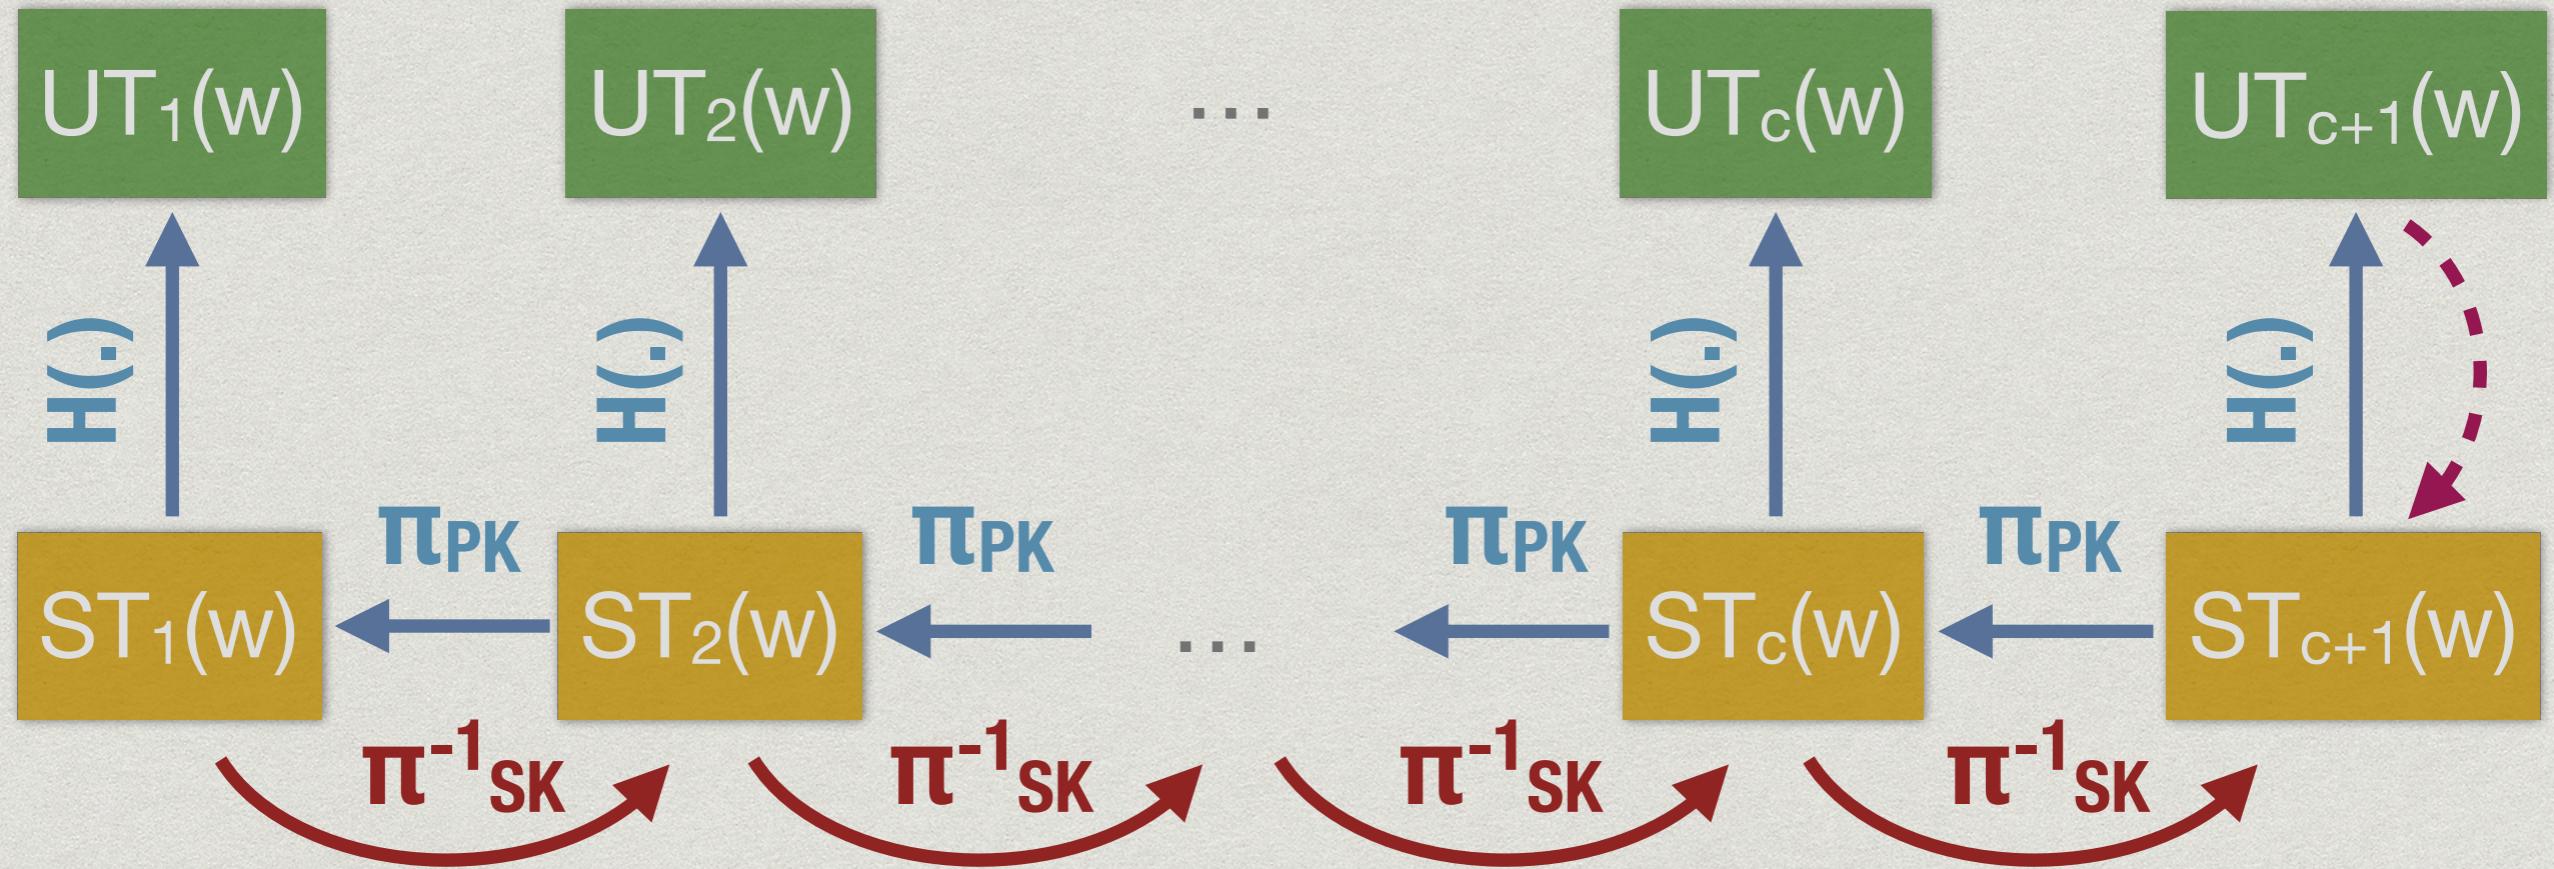




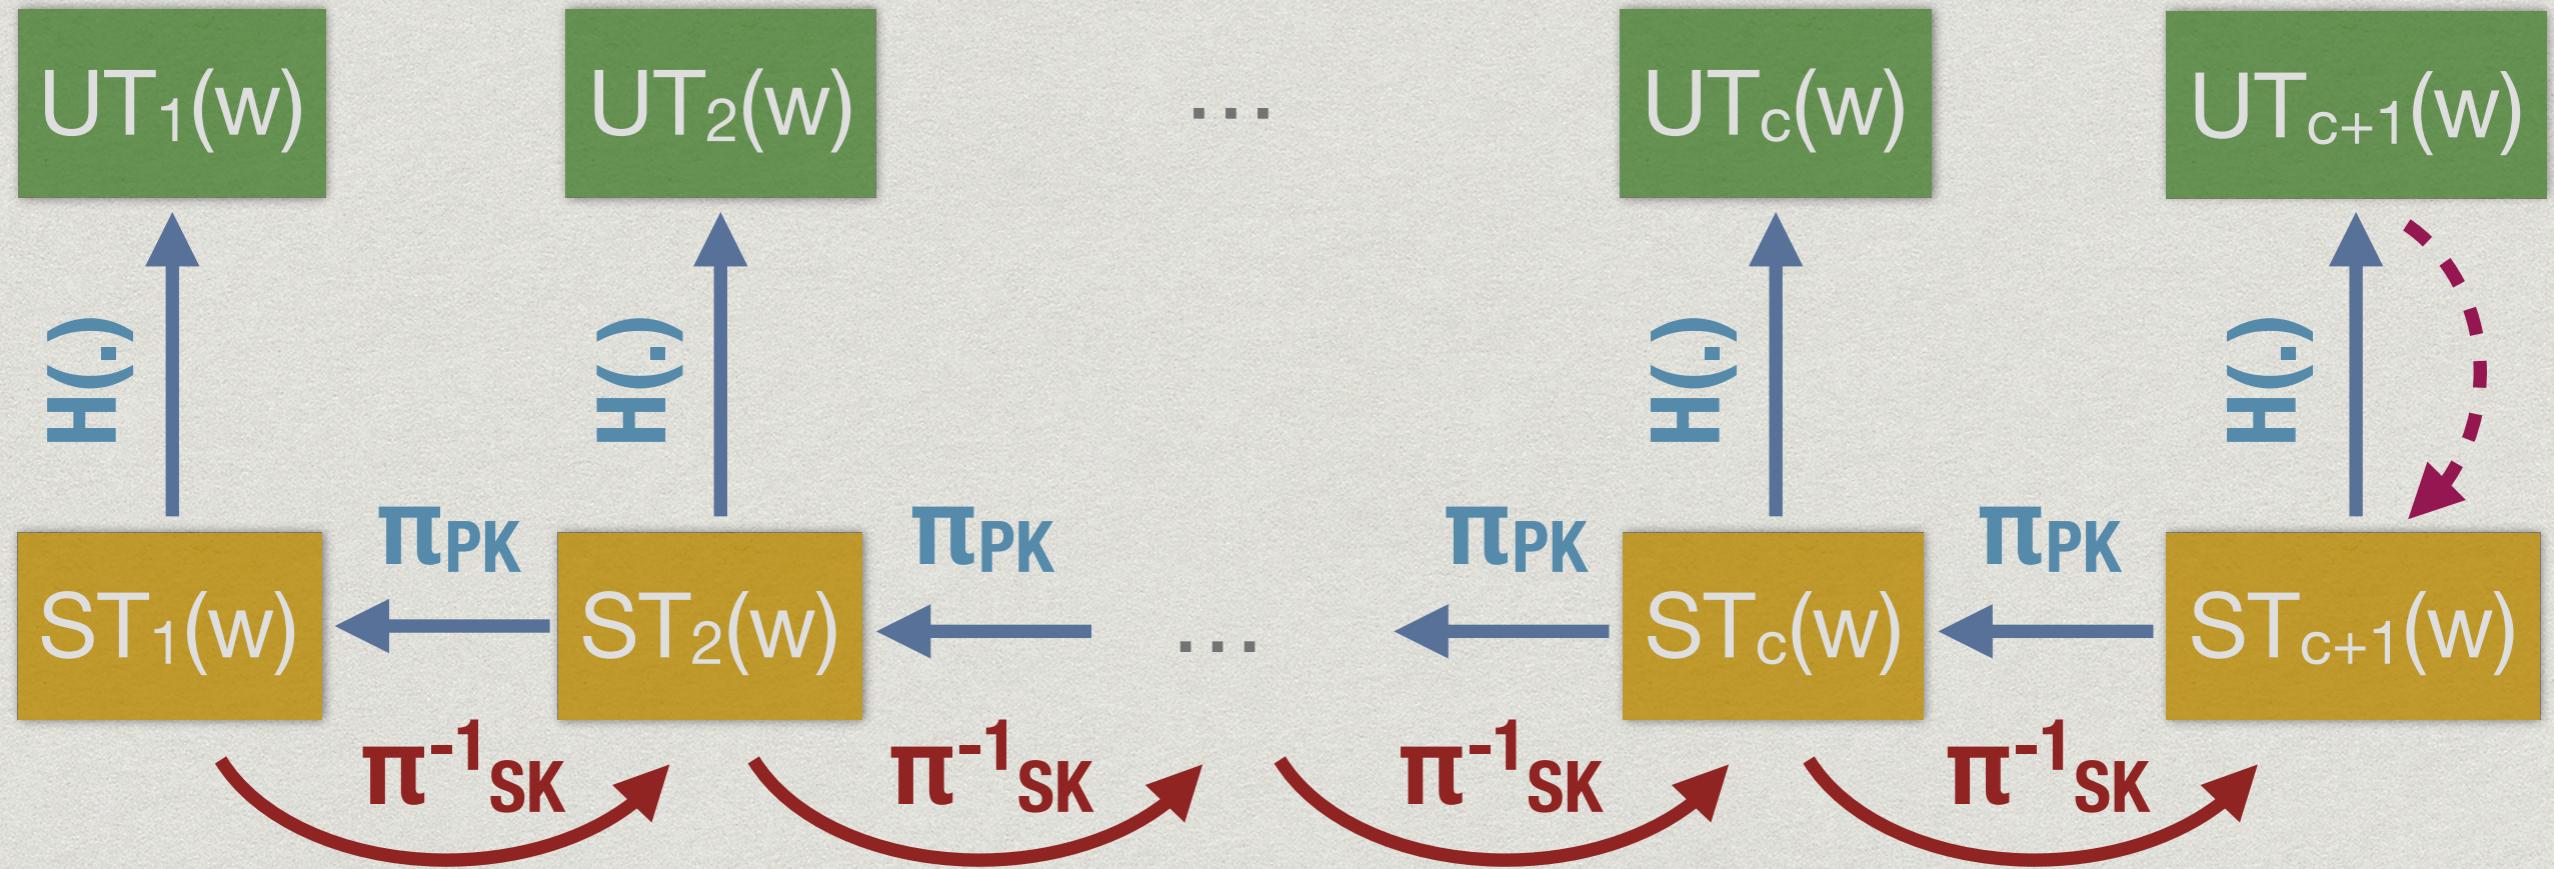
- * Client stores $W[w] := ST_c(w)$



- * Client stores $W[w] := ST_c(w)$
- * Search w : send $ST_c(w)$

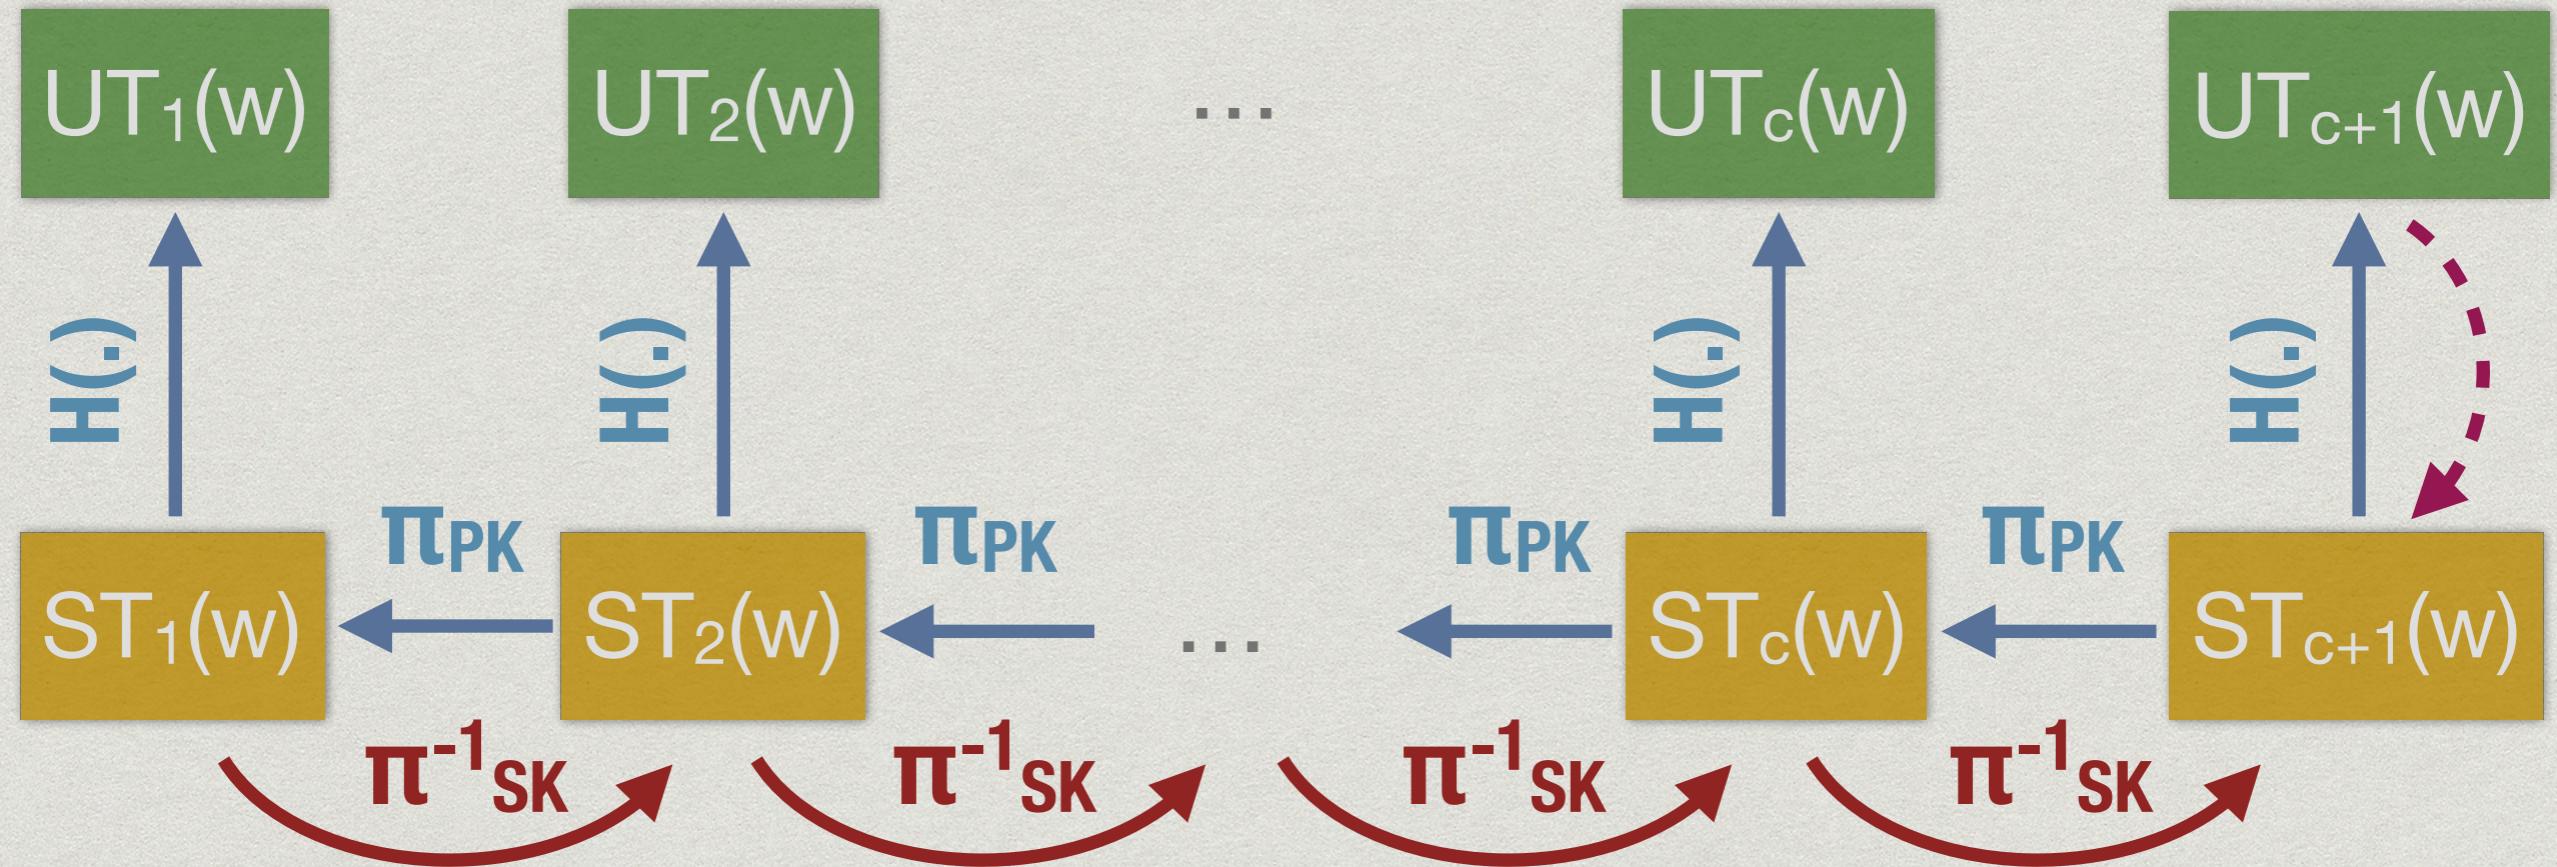


- * Client stores $W[w] := ST_c(w)$
- * Search w : send $ST_c(w)$
- * Update: $W[w] := \pi^{-1}_{SK}(ST_c(w))$



Search:

- * Client: constant
- * Server: $O(|DB(w)|)$

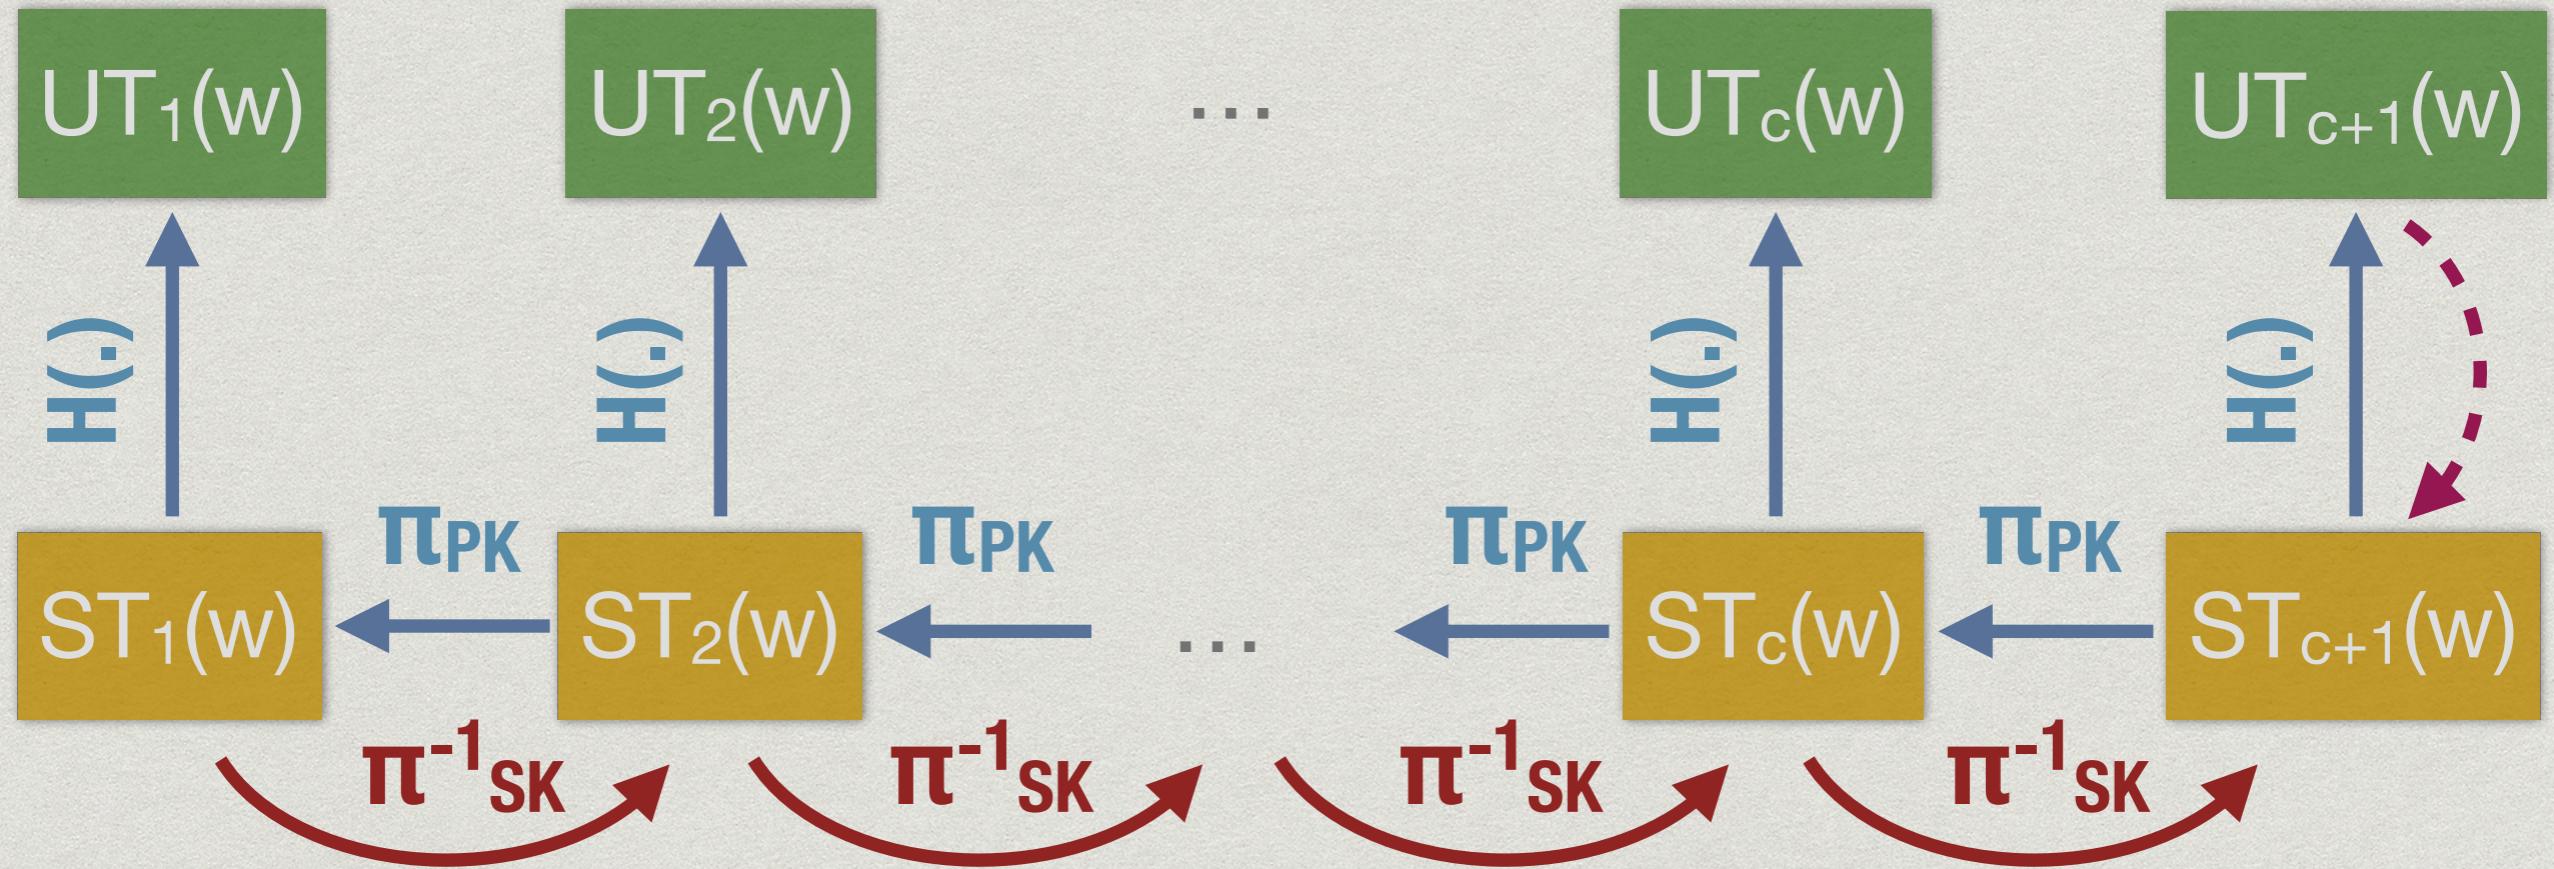


Search:

- * Client: constant
- * Server: $O(|DB(w)|)$

Update:

- * Client: constant
- * Server: constant



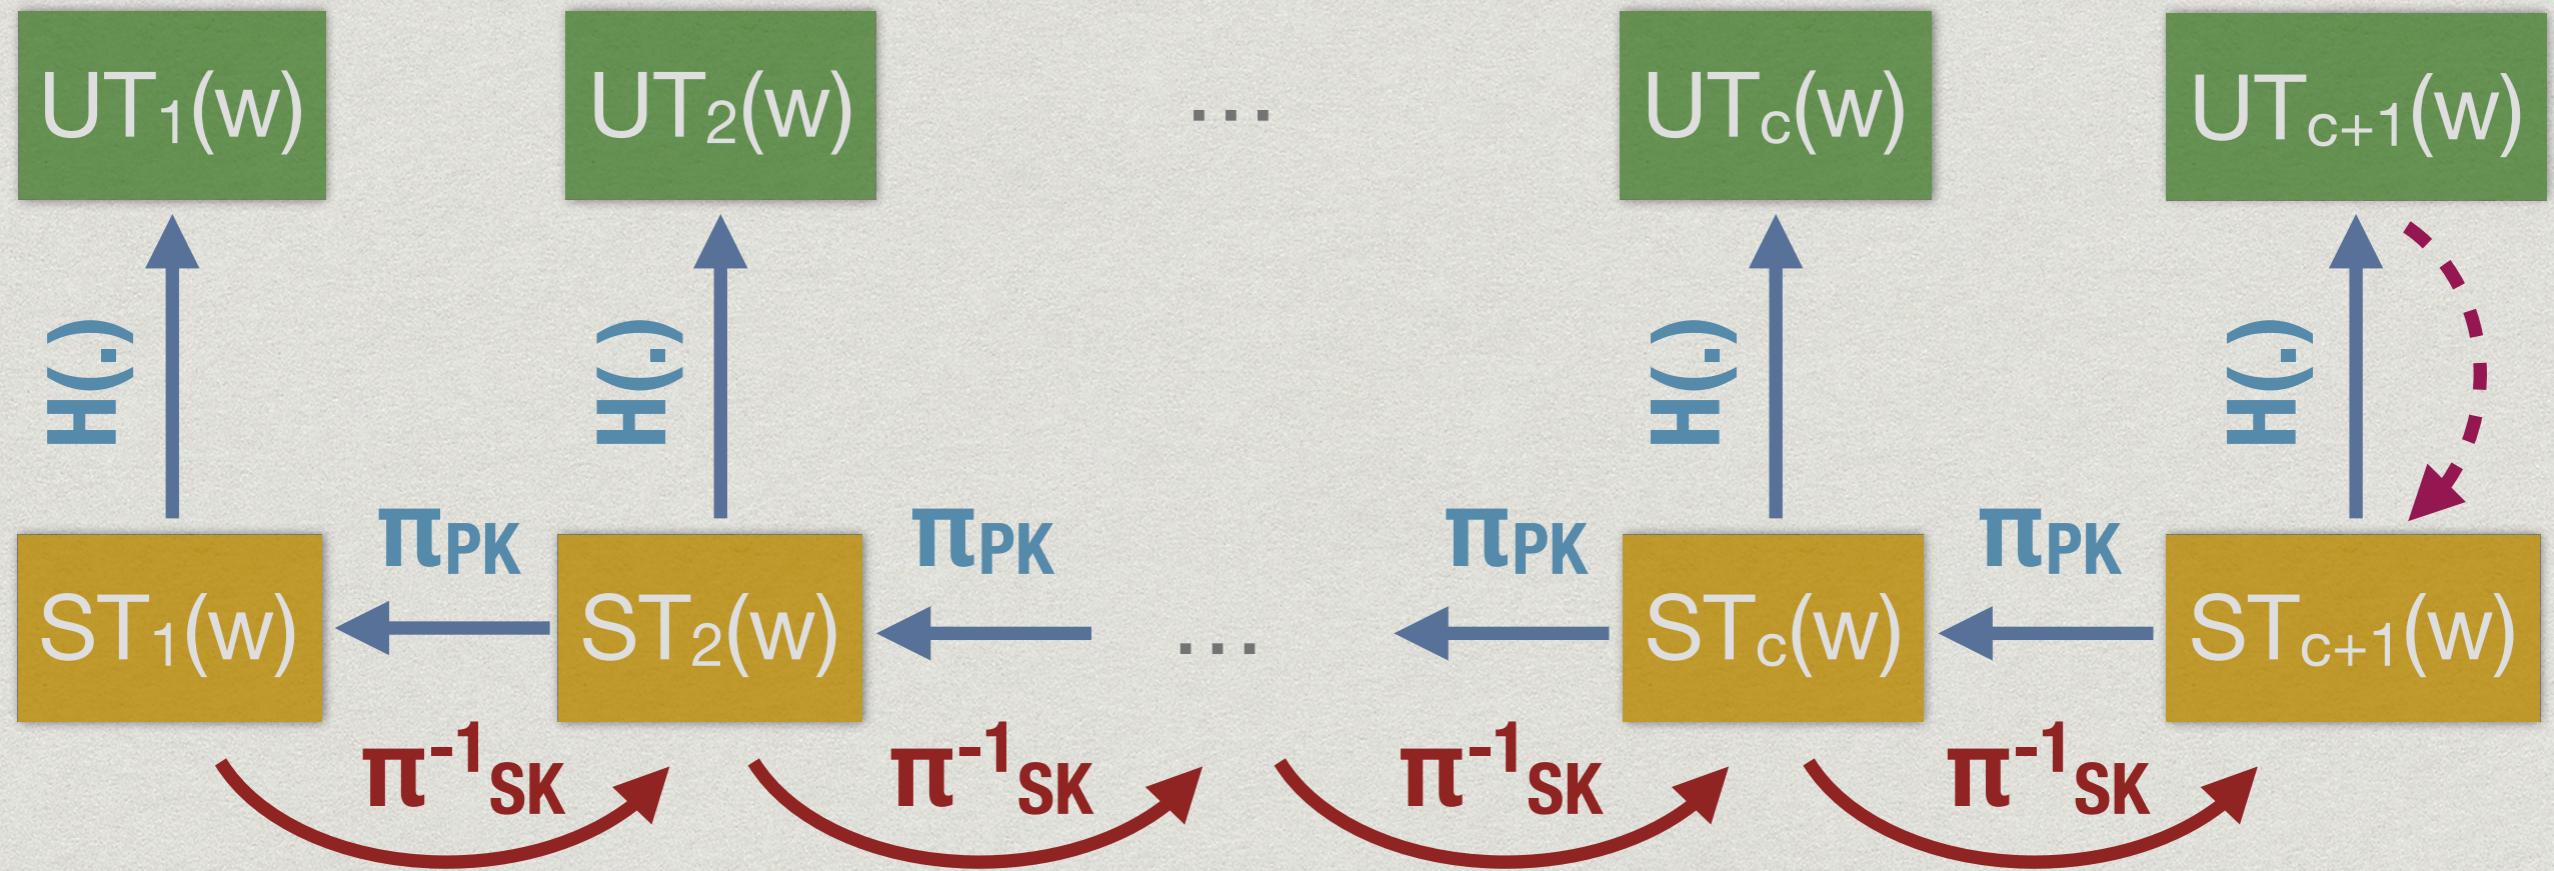
Search:

- * Client: constant
- * Server: $O(|DB(w)|)$

Update:

- * Client: constant
- * Server: constant

Optimal



Storage:

- * Client: $O(K)$
- * Server: $O(|DB|)$

Σοφος

- * TDP π? RSA or Rabin

Σοφος

- * TDP π? RSA or Rabin
 - ✗ Elements (STs) are large (2048 bits).

Σοφος

- * TDP π? RSA or Rabin
 - ✗ Elements (STs) are large (2048 bits).
 - ✗ Huge client storage.

Σοφος

- * TDP π? RSA or Rabin
 - ✗ Elements (STs) are large (2048 bits).
 - ✗ Huge client storage.
- * Client only stores c , pseudo-randomly generates $ST_1(w)$, computes $ST_c(w)$ on the fly

Σοφος

- * TDP π? RSA or Rabin
 - ✗ Elements (STs) are large (2048 bits).
 - ✗ Huge client storage.
- * Client only stores c , pseudo-randomly generates $ST_1(w)$, computes $ST_c(w)$ on the fly
 - ✓ Efficient (non-iterative) using RSA
- * Search is embarrassingly parallelizable

$$x^{d \cdot \dots \cdot d} = x^{(d^c \bmod \phi(N)) \bmod N}$$

Σοφος - Security

- * Update leakage: nothing **Forward private!**
- * Search leakage:
 - search pattern
 - ‘history’ of w: the timestamped list of updates of keyword w

Adaptive security (ROM)

Σοφος - Evaluation

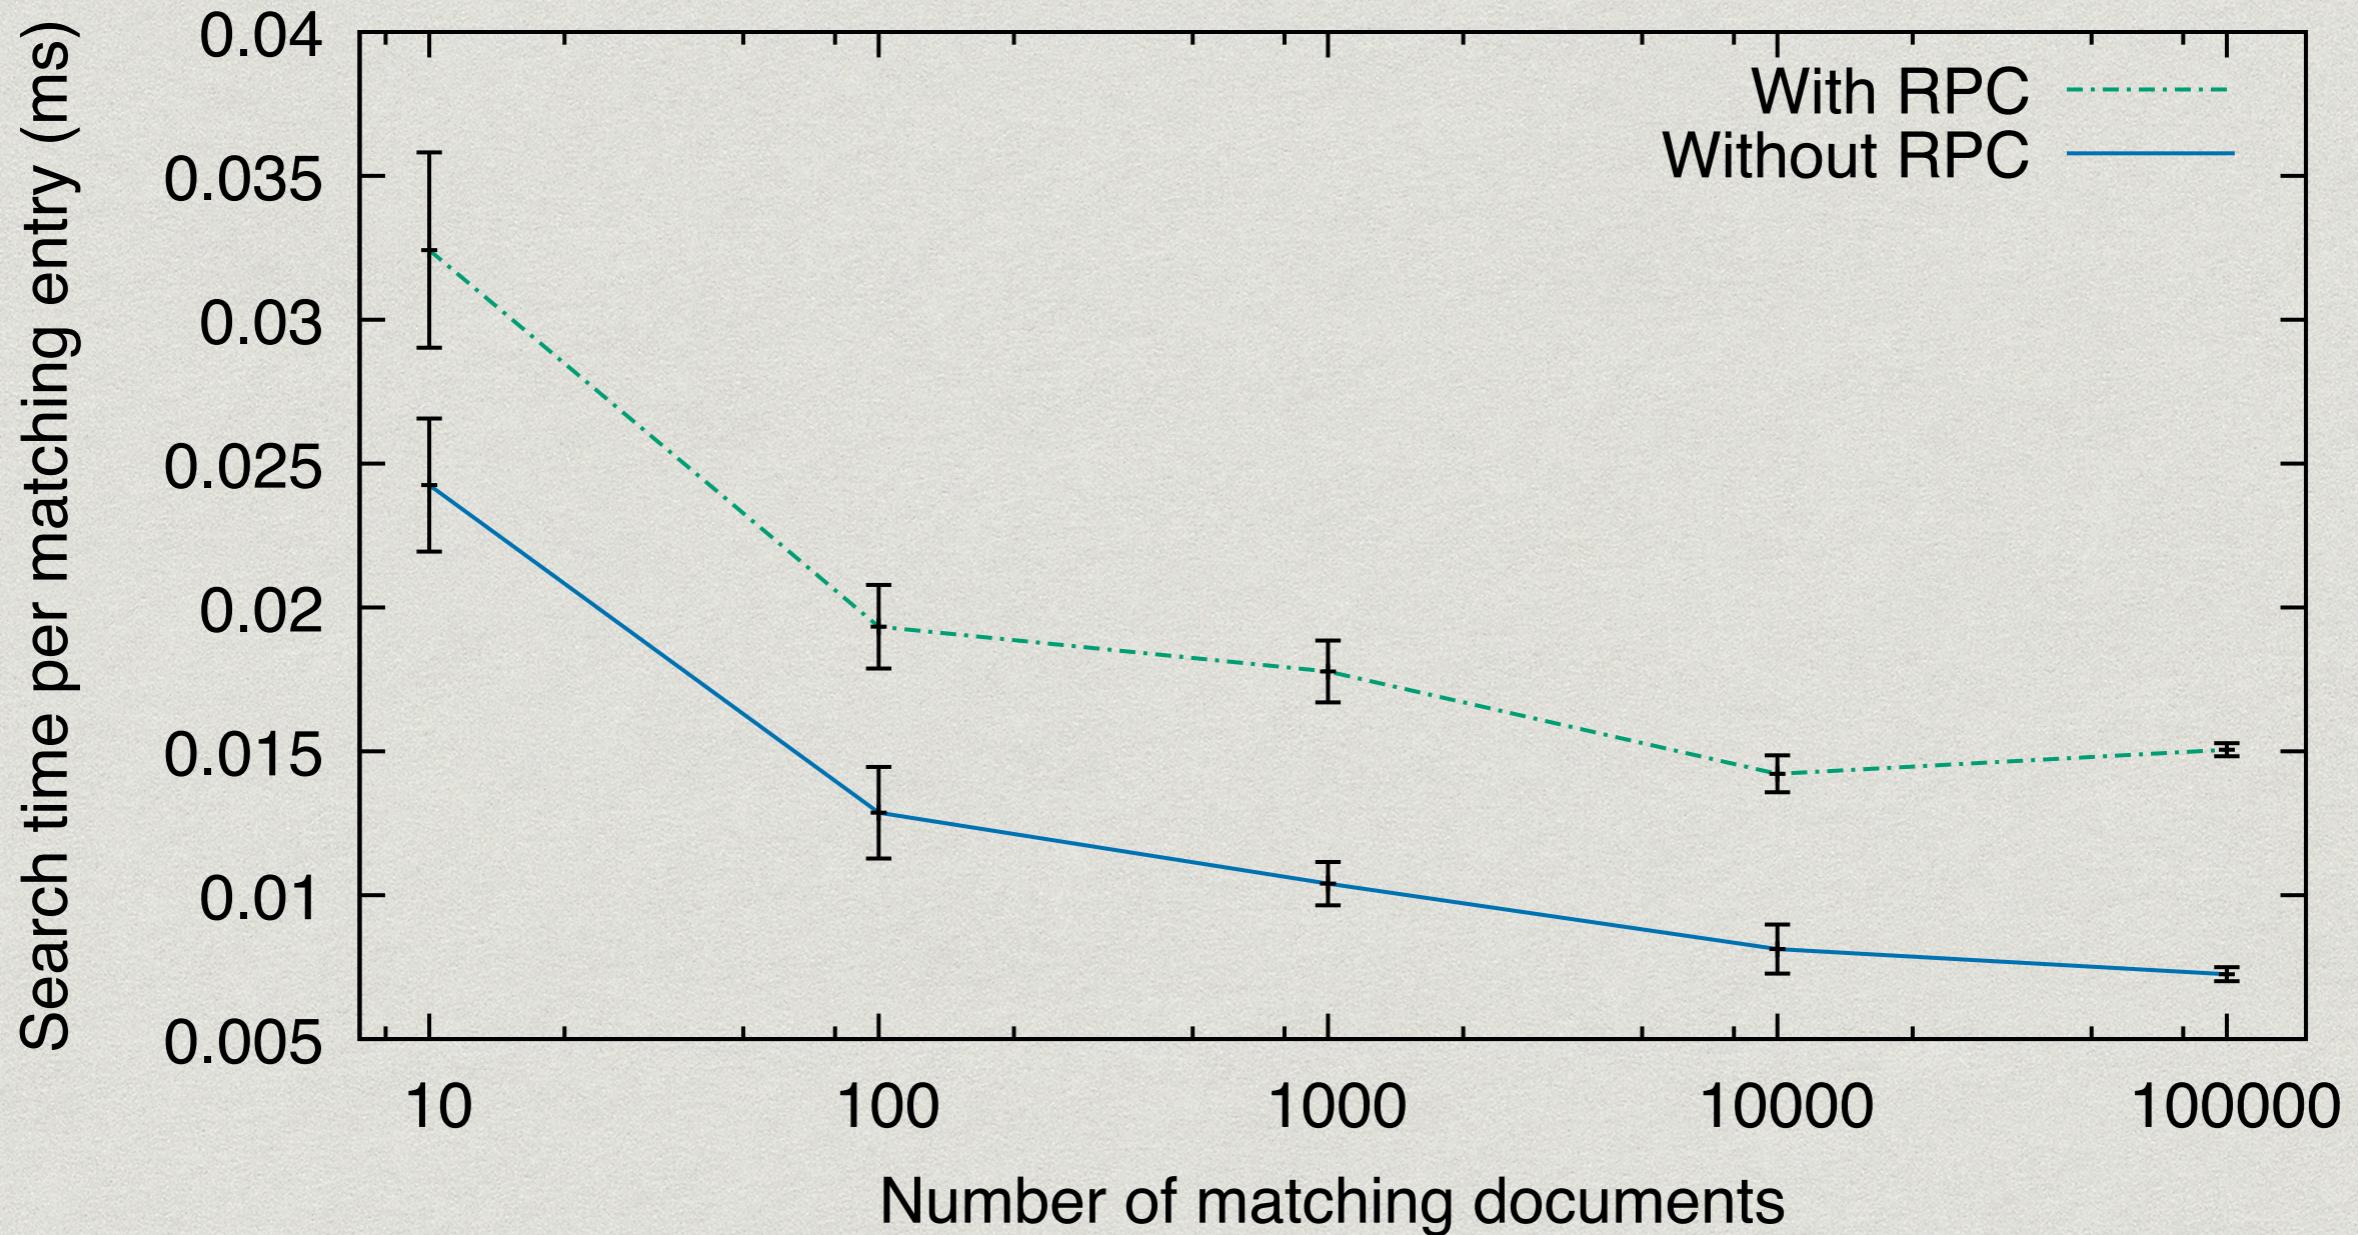
- * C/C++ full fledged implementation
- * Server KVS: RocksDB
- * Evaluated on a desktop computer
4 GHz Core i7 CPU, 8GB RAM, SSD

<https://gitlab.com/sse/sophos>

Σοφος - Evaluation

2M keywords, 140M entries

5.25GB server storage, 64.2 MB Client storage



Σοφος

- * Provable forward privacy

Σοφος

- * Provable forward privacy
- * Very simple

Σοφος

- * Provable forward privacy
- * Very simple
- * Efficient search (IO bounded)

Σοφος

- * Provable forward privacy
- * Very simple
- * Efficient search (IO bounded)
- * Asymptotically efficient update (optimal)

Σοφος

- * Provable forward privacy
- * Very simple
- * Efficient search (IO bounded)
- * Asymptotically efficient update (optimal)
- * In practice, very low update throughput
(4300 p/s - 20x slower than other work)

THANKS!

